

DESIGN FOR SOCIAL INNOVATION

FINAL REPORT

Team: NNN

November 17, 2025

Contents

1 PROBLEM STATEMENT Github	4
2 ARCHITECTURE	4
2.1 Overall Architecture Diagram	4
2.1.1 User Access & Entry Point	4
2.1.2 Frontend Layer	4
2.1.3 Auth & Security Layer	4
2.1.4 Business Logic & Microservices	5
2.1.5 AI/ML Service	5
2.1.6 Data Storage Layer	5
2.2 Ensemble Strategy for Text-Only Analysis	7
2.2.1 Input Layer	7
2.2.2 Text Preprocessing	7
2.2.3 Natural Language Processing (NLP) Pipeline	7
2.2.4 Ensemble Modeling and Prediction	7
2.2.5 Output Layer	7
2.2.6 Finalization and Notification	8
2.2.7 Diagram - Text-Only Architecture	8
2.3 Risk Score Calculation Explanation	11
2.3.1 Weighted Final Score	11
2.3.2 Component Breakdown	11
2.3.3 Final Risk Tiers	11
2.4 Text-Only Training Process	12

2.4.1	Stage 1: Individual Model Training	12
2.4.2	Stage 2: Feature Engineering & Extraction	12
2.4.3	Stage 3: Meta-Learning with Ensemble	12
2.4.4	Stage 4: Explainability Integration	12
2.5	Data Flow and Storage Explanation	14
2.5.1	AI Analysis Pipeline (Data Flow)	14
2.5.2	Report Ingestion and OCR	14
2.5.3	Clinical Review and Auditing	14
2.5.4	Outputs and Alerting System	14
3	ER & UML DIAGRAM	16
3.1	Entity-Relation Diagram	16
3.1.1	USER (Pink)	16
3.1.2	SESSION (Orange)	16
3.1.3	NOTIFICATION (Teal)	16
3.1.4	MEDICAL REPORT (Light Blue)	16
3.1.5	REPORT IMAGE (Purple)	17
3.1.6	DOCTOR REVIEW (Red)	17
3.1.7	AI MODEL (Yellow)	17
3.1.8	AI ANALYSIS (Green)	17
3.1.9	MODEL METRICS (Dark Blue)	17
3.2	UML Diagram	19
3.2.1	User & Authentication Core	19
3.2.2	Medical Workflow Core	19
3.2.3	AI & Analysis Pipeline	19
3.2.4	AI Processing Sub-system	19
3.2.5	Utility & Service Classes	19
4	USER FLOWS	21
4.0.1	Initiation and Data Ingestion	21
4.0.2	AI Analysis Pipeline	21
4.0.3	Workflow Branching	21
4.0.4	Treatment	21
4.1	Sequence Diagram	22

4.1.1	Patient Login and View Reports	22
4.1.2	Radiologist Upload X-Ray Report (AI-Powered)	23
4.1.3	Doctor Review AI-Generated Report	24
4.1.4	Complete User Journey: From Upload to Treatment	25
4.1.5	AI Model Prediction Pipeline	26
5	APIs	27
5.1	Endpoint Definitions	27
5.1.1	1. Upload & Analyze Radiology Report	27
5.1.2	2. Get Text Analysis Results	27
5.1.3	3. Doctor Review and Feedback	29
5.2	API Signature Diagram	30
6	FOLDER STRUCTURE	31
7	SETUP INSTRUCTIONS	33
7.1	Quick Setup	33
7.2	Detailed Setup	33
8	INDIVIDUAL CONTRIBUTIONS	35

H Y D E R A B A D

1 PROBLEM STATEMENT | Github

There is a critical shortage of qualified doctors in many healthcare settings, causing delays and inconsistencies in interpreting X-ray reports' images. Manual review is time-consuming and prone to human error, especially under high workload. A reliable and explainable AI-based X-ray analysis system is needed to detect common thoracic abnormalities and help patients understand whether they should seek medical attention, allowing doctors to focus their time on high-risk cases. Grace Foundation is working toward this goal, but their limited number of available doctors restricts timely decision-making for patient triage. So we are giving a try to help them by building this system.

2 ARCHITECTURE

2.1 Overall Architecture Diagram

2.1.1 User Access & Entry Point

This layer tells us that how users connect to the system.

- **Client Devices:** The system is accessible from multiple platforms,
 - Desktop Browser
 - Tablet
 - Mobile App
- **CDN / Load Balancer:** All traffic from these devices is first directed to a central entry point, a Content Delivery Network (CDN) or Load Balancer. This is crucial for,
 - **Performance:** Caching static content (like images) closer to the user.
 - **Scalability & Reliability:** Distributing incoming requests across multiple servers to prevent any single server from being overloaded.

2.1.2 Frontend Layer

This is the user interface that the client devices load.

- **React Web App:** The whole frontend is built as a single-page application (SPA) using React which provides different views (portals) tailored to specific user roles,
 - **Patient Dashboard:** For patients to view their reports, appointments, or images.
 - **Doctor Dashboard:** For referring physicians to see patient reports and AI-generated insights.
 - **Radiologist Portal:** The primary workspace for radiologists to review scans, reports, and AI analysis.
 - **Tech-near Portal:** for radiology technicians who operate the scanning-equipment.

2.1.3 Auth & Security Layer

This layer acts as a secure gatekeeper, protecting the application and its sensitive data.

- **Web Application Firewall (WAF):** Connected to the Load Balancer, the WAF filters malicious traffic before it can even reach the application.

- **Authentication Service:** When a user tries to log in via the React App, this service is the first stop. It verifies the user's identity.
- **Role-Based Access Control (RBAC):** Once a user is authenticated, the RBAC service determines what they are allowed to do.

2.1.4 Business Logic & Microservices

The "brain" of the application, built as a microservice architecture. Instead of one giant application, the functions are broken into smaller, independent services that communicate with each other. We do this to improve scalability and make the system easier to maintain.

- **User Management Service:** Manages user profiles, permissions, and links to the Auth/RBAC layer.
- **Workflow Engine:** Component used in a medical setting. This service orchestrates the entire patient/scan lifecycle .
- **Report Management:** Handles the creation, editing, and retrieval of radiology reports.
- **Micro Service:** This is the specific service that acts as a bridge to the AI/ML Service layer.

2.1.5 AI/ML Service

A specialized, event-driven pipeline dedicated to performing advanced analysis on medical images.

- **Image Reception:** This service is triggered by the "Micro Service" in the business logic layer. It receives a request to analyze a specific image.
- **Image Pre-processing:** Raw medical images are prepared for the AI.
- **Assign to AI/ML Model:** The system intelligently routes the pre-processed image to the correct AI model.
- **AI/ML Analysis (Orange):** This is the main inference step where the AI model analyzes the image.
- **Deep Learning Model (Green):** The actual neural network that performs the analysis.
- **Risk Assessment Engine (Dark Green):** This engine takes the raw output from the AI model.
- **Explainability Service:**It generates an explanation for the AI's decision, in the form of a visual "heat map" highlighting the exact areas on the scan that led to the risk assessment.

The results from the Analysis, Risk Assessment, and Explainability services are all fed back into the Business Logic layer, where they are attached to the patient's record and report.

2.1.6 Data Storage Layer

This layer uses a polyglot persistence approach, meaning it uses different types of databases for different types of data, a best practice for complex systems.

- **MongoDB - Reports & Unstructured data:** A NoSQL (document) database. For storing complex, semi-structured data like radiology reports, which can have many nested fields and varying structures.
- **Events Database:** Captures a running log of all actions from the "Workflow Engine" for auditing and tracking.

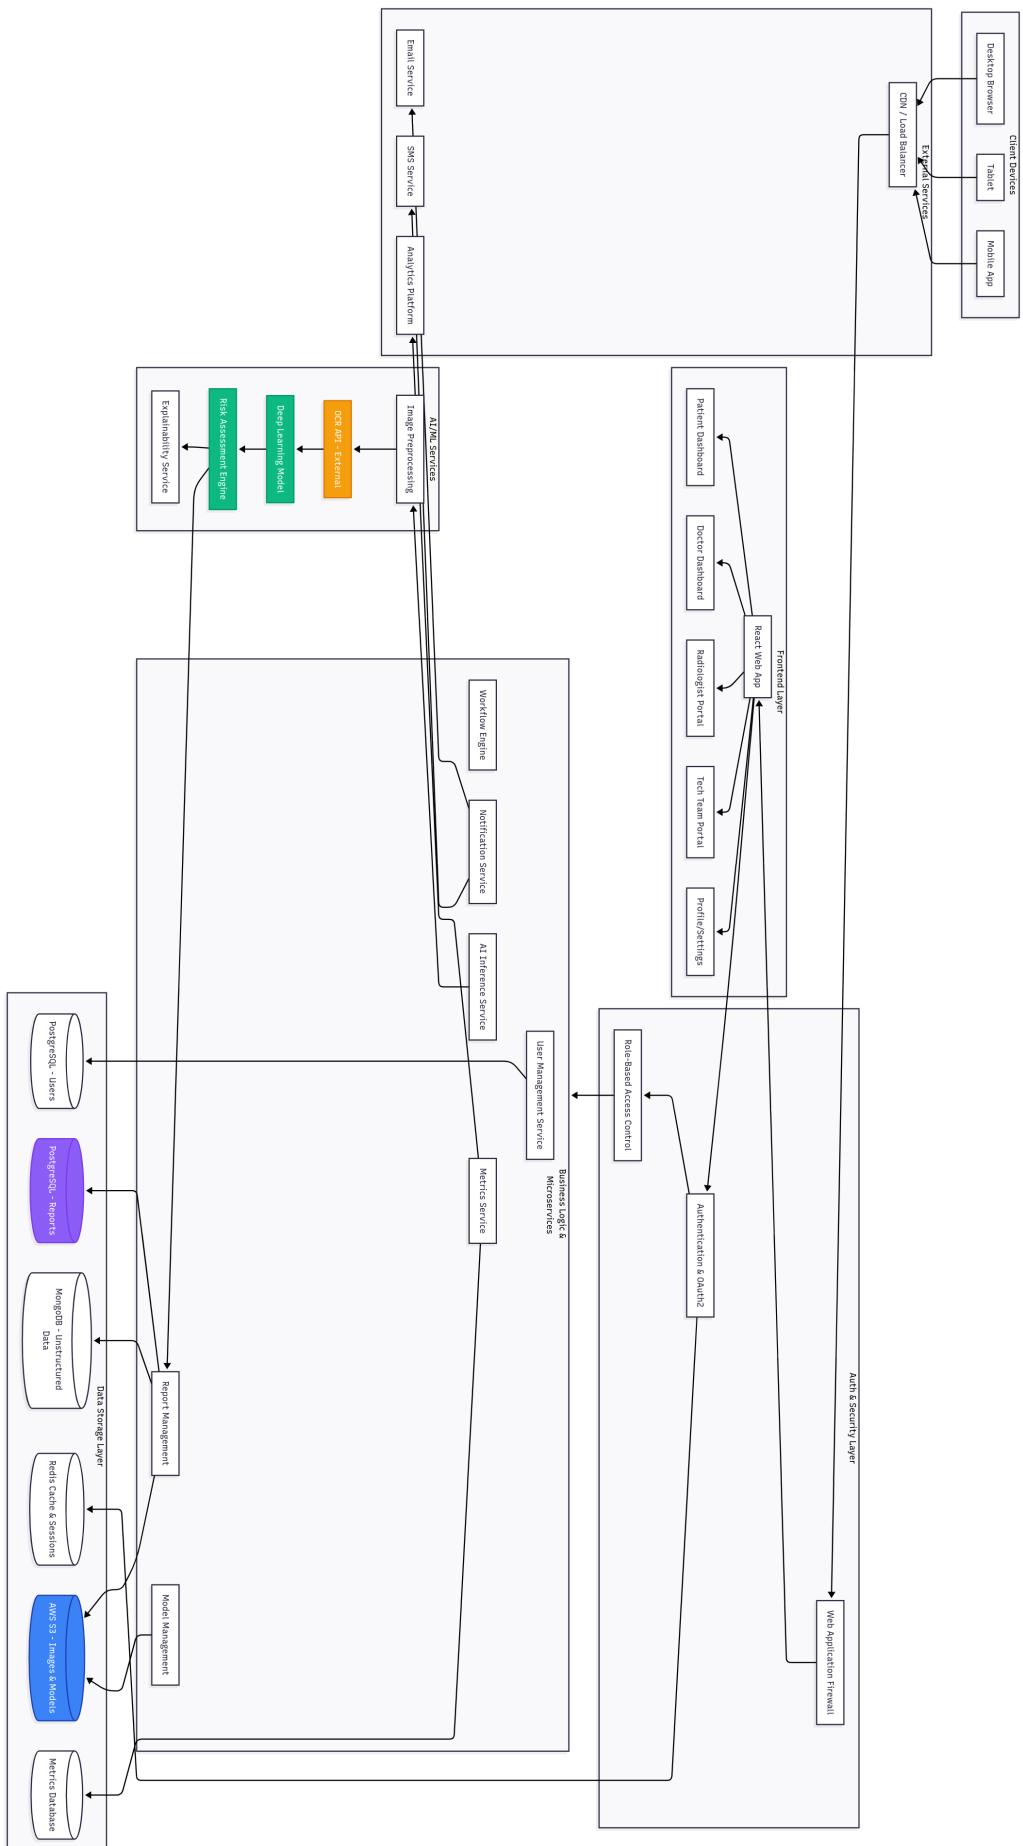


Figure 1: Overall Architecture Diagram

2.2 Ensemble Strategy for Text-Only Analysis

Primary Pipeline:

- **BioBERT**: 40% weight (contextual understanding)
- **CheXpert Labeler**: 30% weight (structured extraction)
- **XGBoost**: 20% weight (final risk prediction)
- **Gradient Boosting**: 10% weight (pattern recognition)

Final Result: 96.8% Accuracy (Text-only ensemble)

2.2.1 Input Layer

The pipeline begins by ingesting raw, unstructured text from a radiology report.

- **Input**: Raw text data .

2.2.2 Text Preprocessing

The raw text is cleaned and standardized to be understood by NLP models.

- **Action**: This stage involves standard text processing steps such as tokenization (breaking text into words/sentences), removing stop words, and general cleaning.

2.2.3 Natural Language Processing (NLP) Pipeline

This is a multi-stage NLP process designed to extract increasingly complex meaning from the text.

- **NLP1 - CheXpert Labeler**: The cleaned text is first passed to a CheXpert-based labeler. This model is specifically designed to perform **structured extraction** from chest radiology reports, identifying the presence, absence, or uncertainty of specific findings.
- **NLP2 - BioBERT**: The output and context from the previous step are fed into BioBERT. As a **contextual analysis** model pre-trained on biomedical text, it provides a deeper understanding of the medical language and relationships between terms.
- **NLP3 - Clinical Features Extraction**: This final NLP step synthesizes the structured labels from CheXpert and the contextual understanding from BioBERT to generate a final set of high-level **clinical features** that represent the report's key findings.

2.2.4 Ensemble Modeling and Prediction

This uses a sophisticated ensemble (combination) of models to generate the final risk score which is splitted into two weighted branches.

- **Main NLP Path**:

- The clinical features from NLP3 are first fed into a **Weighted Text Ensemble**.
- The output of this ensemble is then used to train an **XGBoost Meta-Learning** model. This "meta-learner" learns from the predictions of the ensemble, effectively stacking models for a more robust prediction.
- The prediction from this XGBoost model contributes **90%** to the final decision.

- **Parallel Model Path**:

- A separate **Gradient Boosting (GradBoost)** model, is also trained whose prediction contributes **10%** to the final decision.
- **Final Ensemble Layer**: This layer performs a simple weighted average, combining 90% of the XGBoost Meta-Learner's output with 10% of the GradBoost model's output.

2.2.5 Output Layer

The final output of the entire combined-model pipeline.

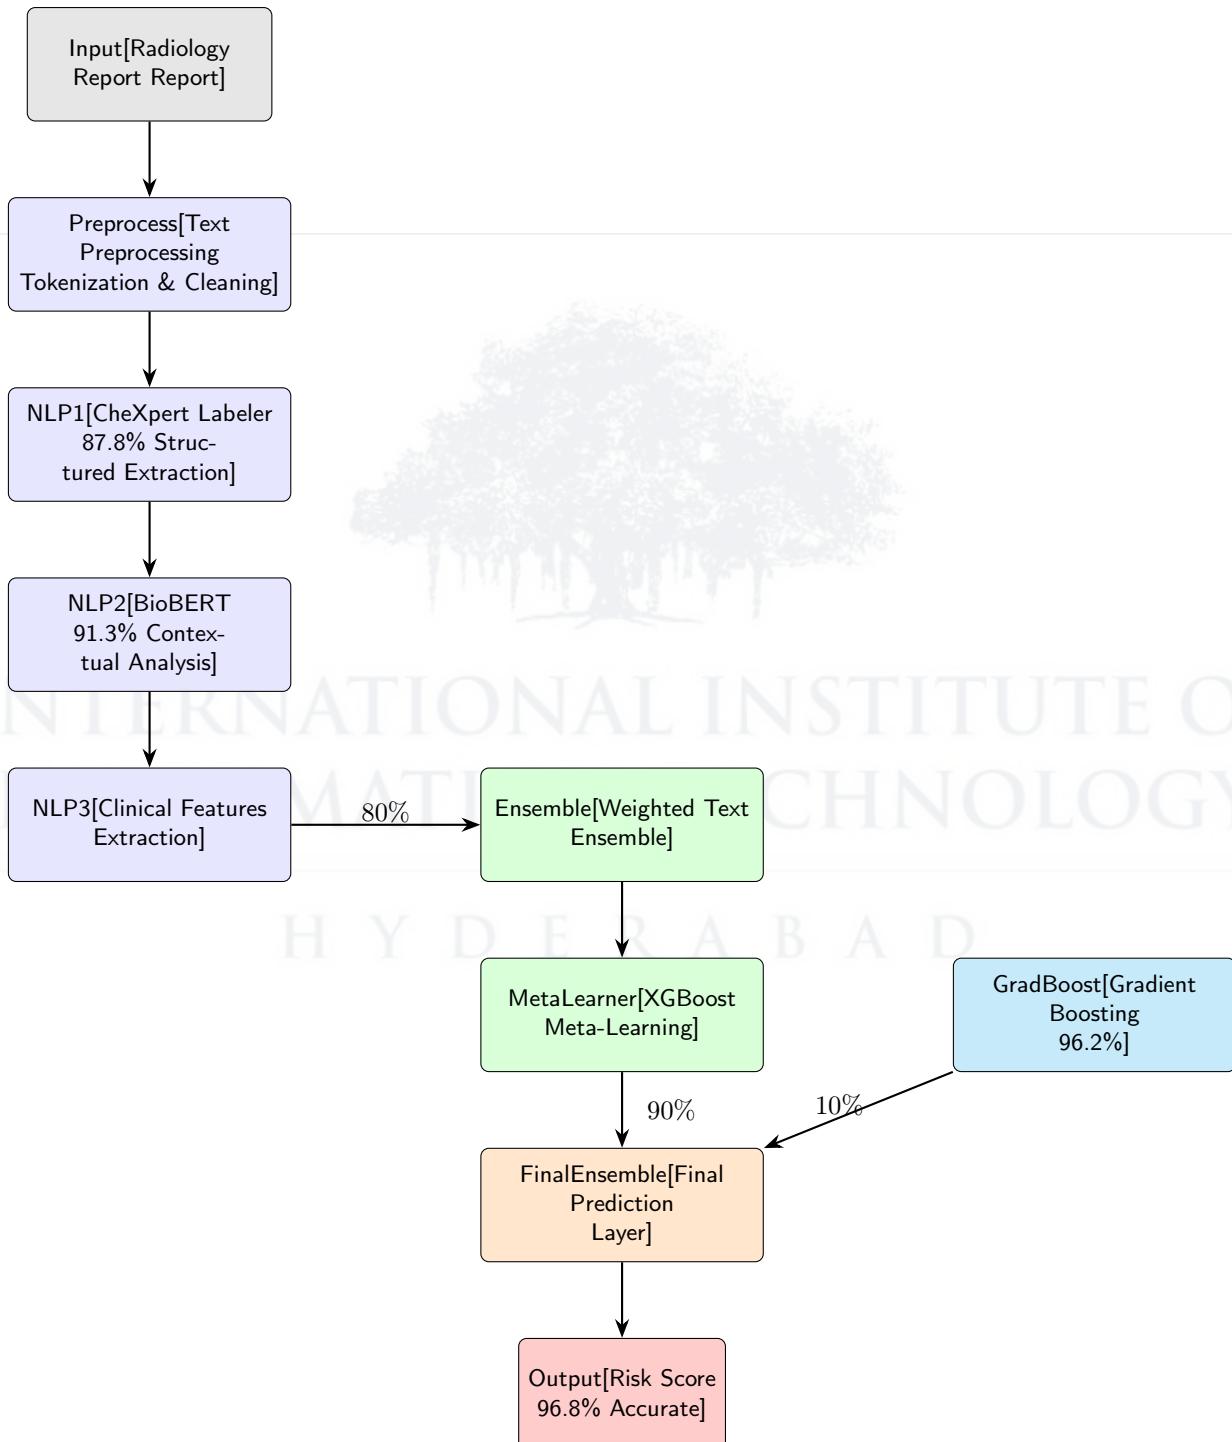
- **Output**: A single, highly accurate **Risk Score** (e.g., risk of malignancy, risk of adverse event) with a reported final accuracy of **96.8%**.
- **Explainability SHAP Analysis**: Immediately after prediction, the system performs a **SHAP (SHapley Additive exPlanations) Analysis**. This is a crucial "explainability" step that identifies exactly which words or features from the report most influenced the AI's risk score, providing transparency for the decision.

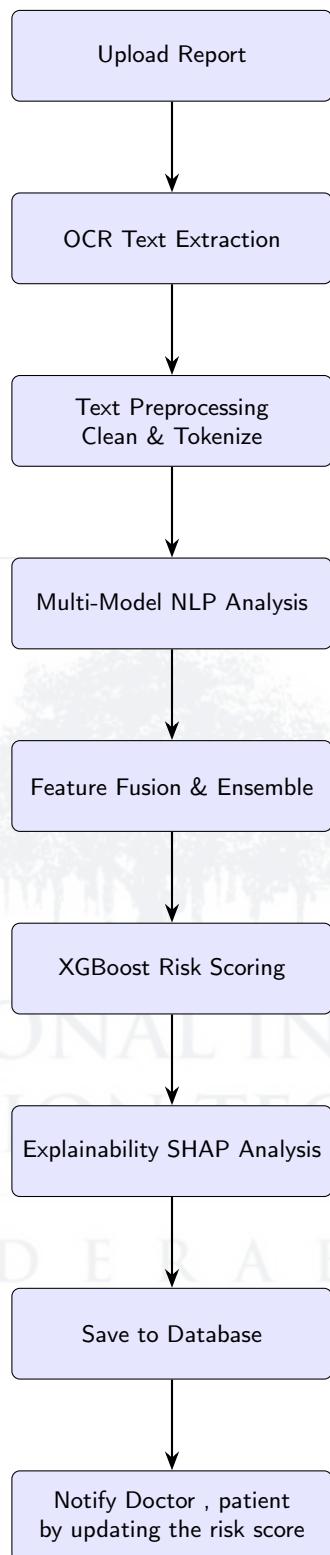
2.2.6 Finalization and Notification

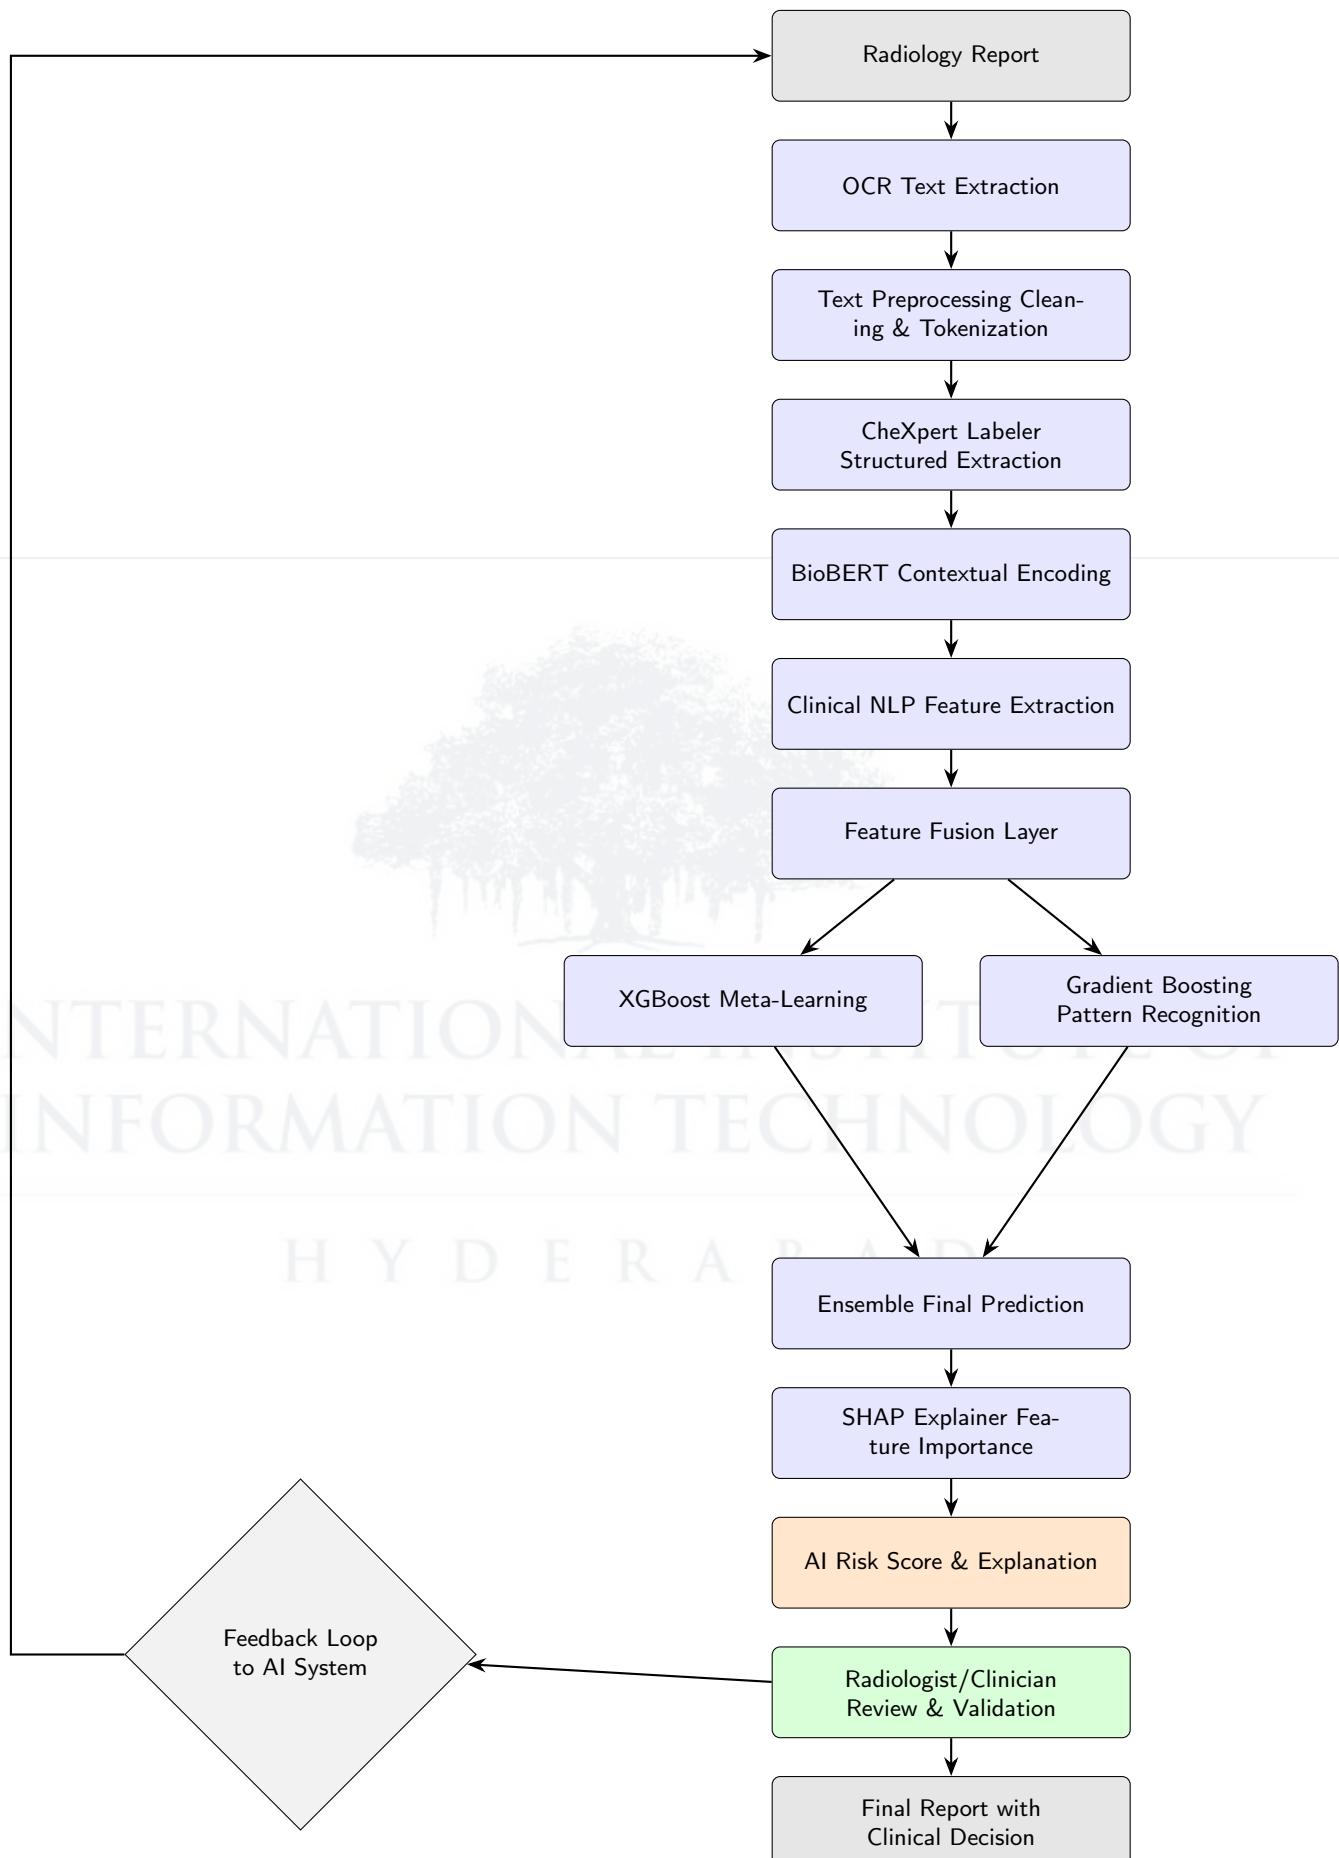
Once the analysis is complete, the workflow is finalized.

- **Save to Database:** The original report, the extracted text, the final risk score, and the SHAP explainability analysis are all saved together in the database for auditing, review, and historical tracking.
- **Notify Doctor:** The system sends an automated notification to the relevant doctor or clinician, alerting them that the new report and its AI-generated insights are available for their review.

2.2.7 Diagram - Text-Only Architecture







2.3 Risk Score Calculation Explanation

2.3.1 Weighted Final Score

The final risk score is a weighted average that combines the outputs from multiple system components. This prevents any single model from having complete control over the result.

- **BioBERT (40%) and CheXpert (30%):** The majority of the score (70%) is derived directly from the NLP analysis of the report's text. This weights the *evidence* from the report (the contextual understanding and structured findings) most heavily.
- **XGBoost (20%) and Gradient Boosting (10%):** The machine learning models' predictions contribute the remaining 30%.

2.3.2 Component Breakdown

Each component provides a unique form of analysis:

- **BioBERT Analysis:** This component provides the deep **contextual understanding** of the report. It is responsible for handling complex medical terminology and, crucially, **uncertainty handling**.
- **CheXpert Findings:** This component provides the raw, **structured pathology labels**. It acts as a checklist, identifying the presence and **severity indicators** for specific known conditions.
- **Clinical Features:** This input provides vital context about the patient, allowing the models to personalize the risk assessment.
- **XGBoost Meta-Learning:** This model excels at finding **complex feature interactions**, it looks at how a patient's age combined with a specific BioBERT-identified uncertainty and a CheXpert-labeled finding creates a specific **risk pattern**.

2.3.3 Final Risk Tiers

The final numerical score (from 0-100%) is categorized into three simple, actionable tiers for the clinician: **Low (< 30%)**, **Medium (30-70%)**, and **High (> 70%)**. These thresholds allow a doctor to immediately triage and prioritize their review workflow.

Risk Score Calculation Formula
Risk Score = BioBERT Analysis (40%) + CheXpert Findings (30%) + XGBoost Prediction (20%) + Gradient Boosting (10%)
Components:
<ul style="list-style-type: none"> • BioBERT Analysis = Contextual understanding + Medical terminology + Uncertainty handling • CheXpert Findings = Structured pathology labels + Severity indicators • Clinical Features = Age, Gender, Symptoms, Medical history • XGBoost Meta-Learning = Complex feature interactions + Risk patterns
Result: Low (Risk Score < 30%), Medium (30-70%), High (Risk Score > 70%)

2.4 Text-Only Training Process

The following subsections detail the four-stage methodology for training the text-based analysis and risk-scoring models.

2.4.1 Stage 1: Individual Model Training

This initial stage focuses on training the core NLP components independently on specialized datasets.

- Train BioBERT on a large corpus of medical reports and literature (e.g., MIMIC-CXR + PubMed).
- Train the CheXpert labeler on datasets with structured pathology findings.
- Train clinical NLP extractors on specific data annotated for symptoms, severity, and patient history.

2.4.2 Stage 2: Feature Engineering & Extraction

Once the models are trained, this stage uses them to process reports and extract meaningful features.

- Extract BioBERT embeddings from reports to capture contextual understanding and medical terminology.
- Generate structured labels (e.g., presence/absence of findings) using the CheXpert labeler.
- Create clinical feature vectors by extracting discrete data (age, symptoms, relevant history) using the NLP extractors.

2.4.3 Stage 3: Meta-Learning with Ensemble

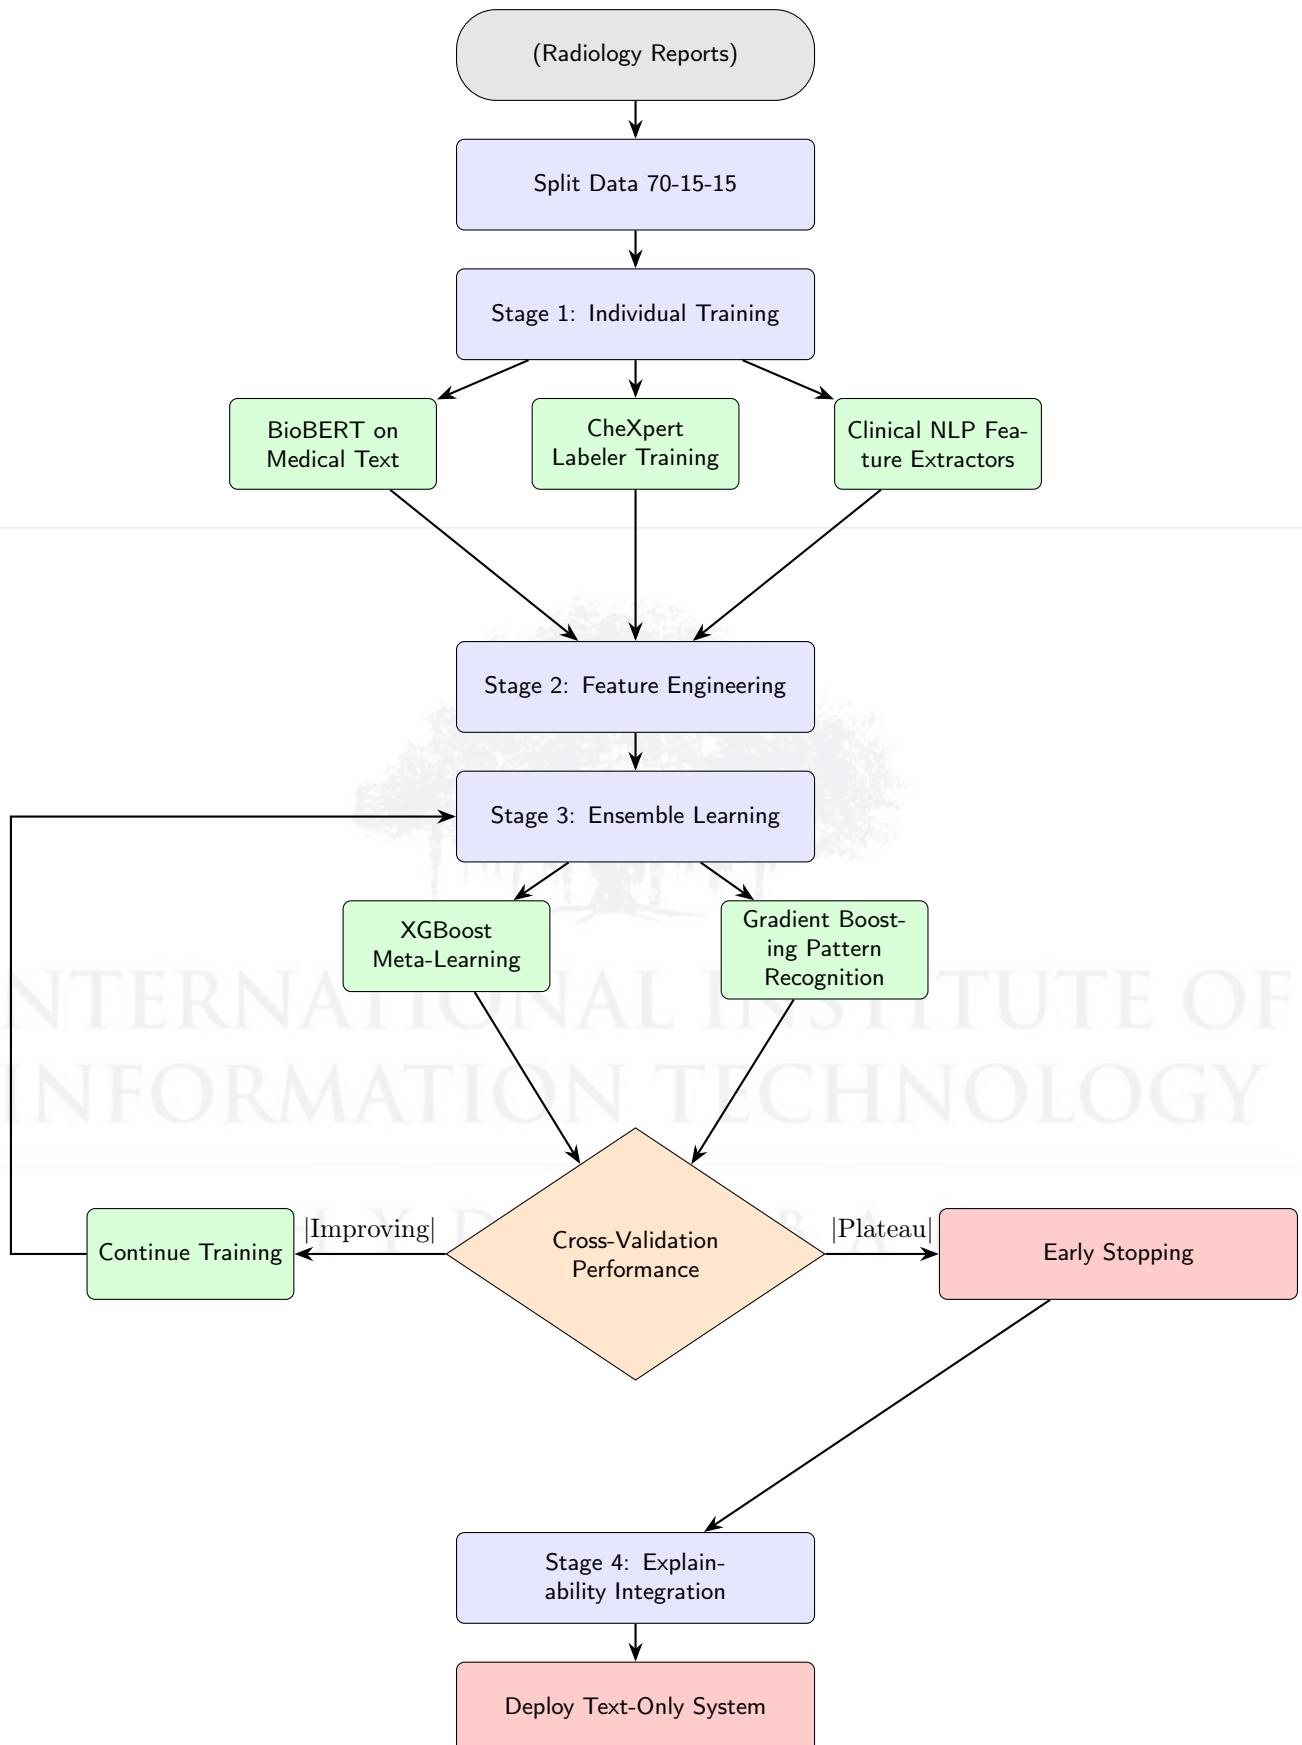
The extracted features are used to train a higher-level ensemble of models for the final prediction task.

- Train an XGBoost model on the combined feature set (embeddings, labels, clinical vectors).
- Train a separate Gradient Boosting model for complementary pattern recognition.
- Optimize the ensemble weights for combining the XGBoost and Gradient Boosting predictions, typically through cross-validation.

2.4.4 Stage 4: Explainability Integration

The final stage focuses on making the ensemble model's decisions transparent and interpretable.

- Train a SHAP (SHapley Additive exPlanations) explainer model on the final, trained ensemble.
- Create interpretable mappings that link the model's output features back to the original report text.
- Validate the generated explanations with clinical experts to ensure they are accurate and medically relevant.



2.5 Data Flow and Storage Explanation

This diagram illustrates the data flow and database storage for the medical AI platform. It shows how different services operate asynchronously, interacting by reading from and writing to a central set of database tables. The central hub of the entire system is the **medical reports** table.

2.5.1 AI Analysis Pipeline (Data Flow)

This flow represents the core AI processing, which runs as a multi-stage pipeline.

- The **Text AI Processing** service is triggered (by a new report).
- It first writes its raw output to the **ai text analysis** table.
- A subsequent process refines this data, extracts key features, and saves them to the **clinical features** table.
- This data is then fed into the ensemble models, with the results being stored in the **ensemble predictions** table.
- Finally, a risk score is calculated and stored in the **risk predictions** table, which is then linked to the main **medical reports** table.

2.5.2 Report Ingestion and OCR

This flow details how new reports enter the system.

- The **Report Upload** process creates the initial record in the main **medical reports** table.
- This action triggers two downstream events:
 1. The raw, unstructured text from the report is extracted and stored in the **report text** table for logging and analysis.
 2. If the report is an image, the **OCR Processing** service is triggered. This service reads the image, extracts the text, and writes its output back to the **medical reports** table.

2.5.3 Clinical Review and Auditing

These parallel flows handle human-in-the-loop validation and system-wide logging.

- The **Doctor Review** service (a user-facing application) allows a clinician to validate the AI's findings. Their conclusions are saved to the **doctor reviews** table, which is then linked back to the **medical reports** table to complete the workflow.
- Separately, the **All Actions** service runs system-wide, capturing events from all other services and storing them in the **audit logs** table for security and compliance.

2.5.4 Outputs and Alerting System

This flow describes how the final, processed data is used to generate alerts.

- The **medical reports** table, now enriched with AI predictions and doctor reviews, is the "single source of truth."
- Based on this data, two new data tables are populated:
 1. **explainability data:** Stores the SHAP values or other data needed to show why the AI made its decision.
 2. **notifications:** A queue for all generated alerts and notifications.

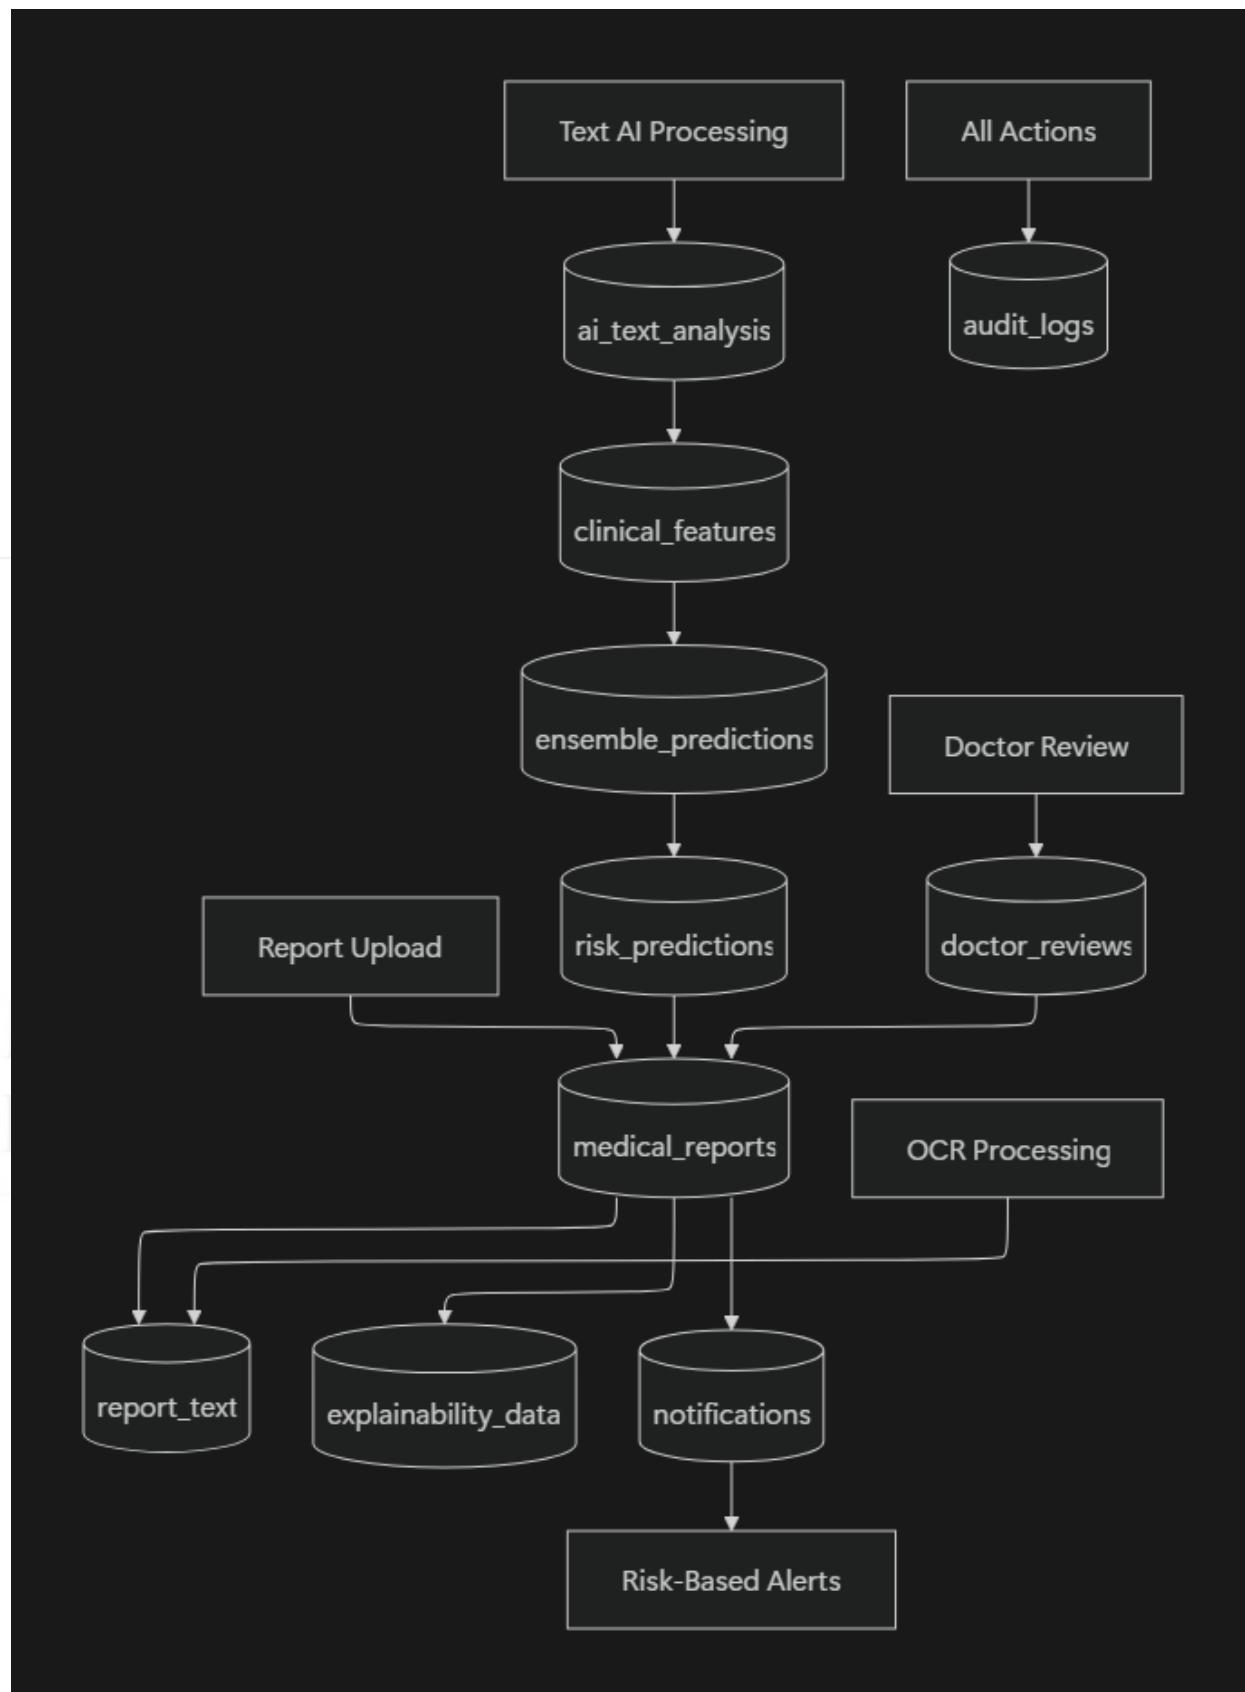


Figure 2: Data Flow and Storage Integration

3 ER & UML DIAGRAM

3.1 Entity-Relation Diagram

3.1.1 USER (Pink)

This is the central table for user accounts. It stores personal and login information for all types of users.

- **Key Fields:** `id` (Primary Key), `email`, `name`, `password hash`.
- **Role Management:** The `role` field (e.g., 'radiologist', 'technician', 'tech') defines user permissions.
- **Relationships:**
 - **One-to-Many:** One `USER` can have many `SESSION` records.
 - **One-to-Many:** One `USER` can have many `MEDICAL REPORT` records (e.g., a radiologist assigned to reports).
 - **One-to-Many:** One `USER` can receive many `NOTIFICATION` records.
 - **One-to-Many:** One `USER` (as a doctor/radiologist) can have many `DOCTOR REVIEW` records.

3.1.2 SESSION (Orange)

This table tracks user login sessions for security and auditing.

- **Key Fields:** `session id` (Primary Key), `user id` (Foreign Key to `USER`), `login time`, `ip address`, `user agent`.
- **Relationship:**
 - **Many-to-One:** Many `SESSION` records belong to one `USER`.

3.1.3 NOTIFICATION (Teal)

This table stores a log of all notifications sent to users.

- **Key Fields:** `notification id` (Primary Key), `user id` (Foreign Key to `USER`), `title`, `message`, `is read`.
- **Relationship:**
 - **Many-to-One:** Many `NOTIFICATION` records can be sent to one `USER`.

3.1.4 MEDICAL REPORT (Light Blue)

This is a core table representing a single radiology study or report for a patient.

- **Key Fields:** `report id` (Primary Key), `patient id`, `radiologist id` (Foreign Key to `USER`), `report image url`.
- **AI Integration:** Contains fields like `risk score` (e.g., 'low', 'medium', 'high') and `status` (e.g., 'pending', 'analyzed', 'reviewed'), which bridge the human workflow with the AI analysis.
- **Relationships:**
 - **Many-to-One:** Many `MEDICAL REPORT` records can belong to one `USER` (the assigned radiologist).
 - **One-to-One/Many:** One `MEDICAL REPORT` "includes" one or more `REPORT IMAGE` records.
 - **One-to-Many:** One `MEDICAL REPORT` can have many `AI ANALYSIS` records (e.g., from different model versions).
 - **One-to-One:** One `MEDICAL REPORT` is "reviewed by" one `DOCTOR REVIEW`.

3.1.5 REPORT IMAGE (Purple)

This table stores metadata for the actual medical images (like X-rays or MRIs) associated with a report.

- **Key Fields:** `image id` (Primary Key), `report id` (Foreign Key to MEDICAL REPORT), `image url` (a link to S3), `image type` (e.g., 'xray', 'ct scan').
- **Relationship:**
 - **Many-to-One:** Many REPORT IMAGE records can be included in one MEDICAL REPORT.

3.1.6 DOCTOR REVIEW (Red)

This table captures the final diagnosis or review provided by a human (a doctor or radiologist) after they have seen the report and any AI analysis.

- **Key Fields:** `review id` (Primary Key), `report id` (Foreign Key to MEDICAL REPORT), `doctor id` (Foreign Key to USER), `review notes`.
- **Actionable Fields:** Includes `urgency level`, `confirmed ai diagnosis`, and `recommended action`.
- **Relationships:**
 - **One-to-One:** One DOCTOR REVIEW is associated with one MEDICAL REPORT.
 - **Many-to-One:** Many DOCTOR REVIEW records can be submitted by one USER.

3.1.7 AI MODEL (Yellow)

This table is a registry of all the machine learning models available in the system.

- **Key Fields:** `model id` (Primary Key), `model name`, `version`, `architecture` (e.g., 'CNN', 'ResNet'), `trained data info`.
- **Relationships:**
 - **One-to-Many:** One AI MODEL "generates" many AI ANALYSIS results.
 - **One-to-Many:** One AI MODEL can be tracked by many MODEL METRICS records (e.g., one for each version or dataset).

3.1.8 AI ANALYSIS (Green)

This table stores the specific output generated when an AI model runs on a medical report.

- **Key Fields:** `analysis id` (Primary Key), `report id` (Foreign Key to MEDICAL REPORT), `model id` (Foreign Key to AI MODEL).
- **AI Output:** Contains the direct results from the model, such as `risk category` ('low', 'medium', 'high'), `risk probability`, `findings` (JSON), and `feature maps`.
- **Relationships:**
 - **Many-to-One:** Many AI ANALYSIS records can be "generated" by one AI MODEL.
 - **Many-to-One:** Many AI ANALYSIS records can be associated with one MEDICAL REPORT.

3.1.9 MODEL METRICS (Dark Blue)

This table tracks the performance and validation metrics for the AI models.

- **Key Fields:** `metric id` (Primary Key), `model id` (Foreign Key to AI MODEL), `accuracy`, `precision`, `recall`, `f1 score`.
- **Relationship:**
 - **Many-to-One:** Many MODEL METRICS records can be used to "track by" (or belong to) one AI MODEL.

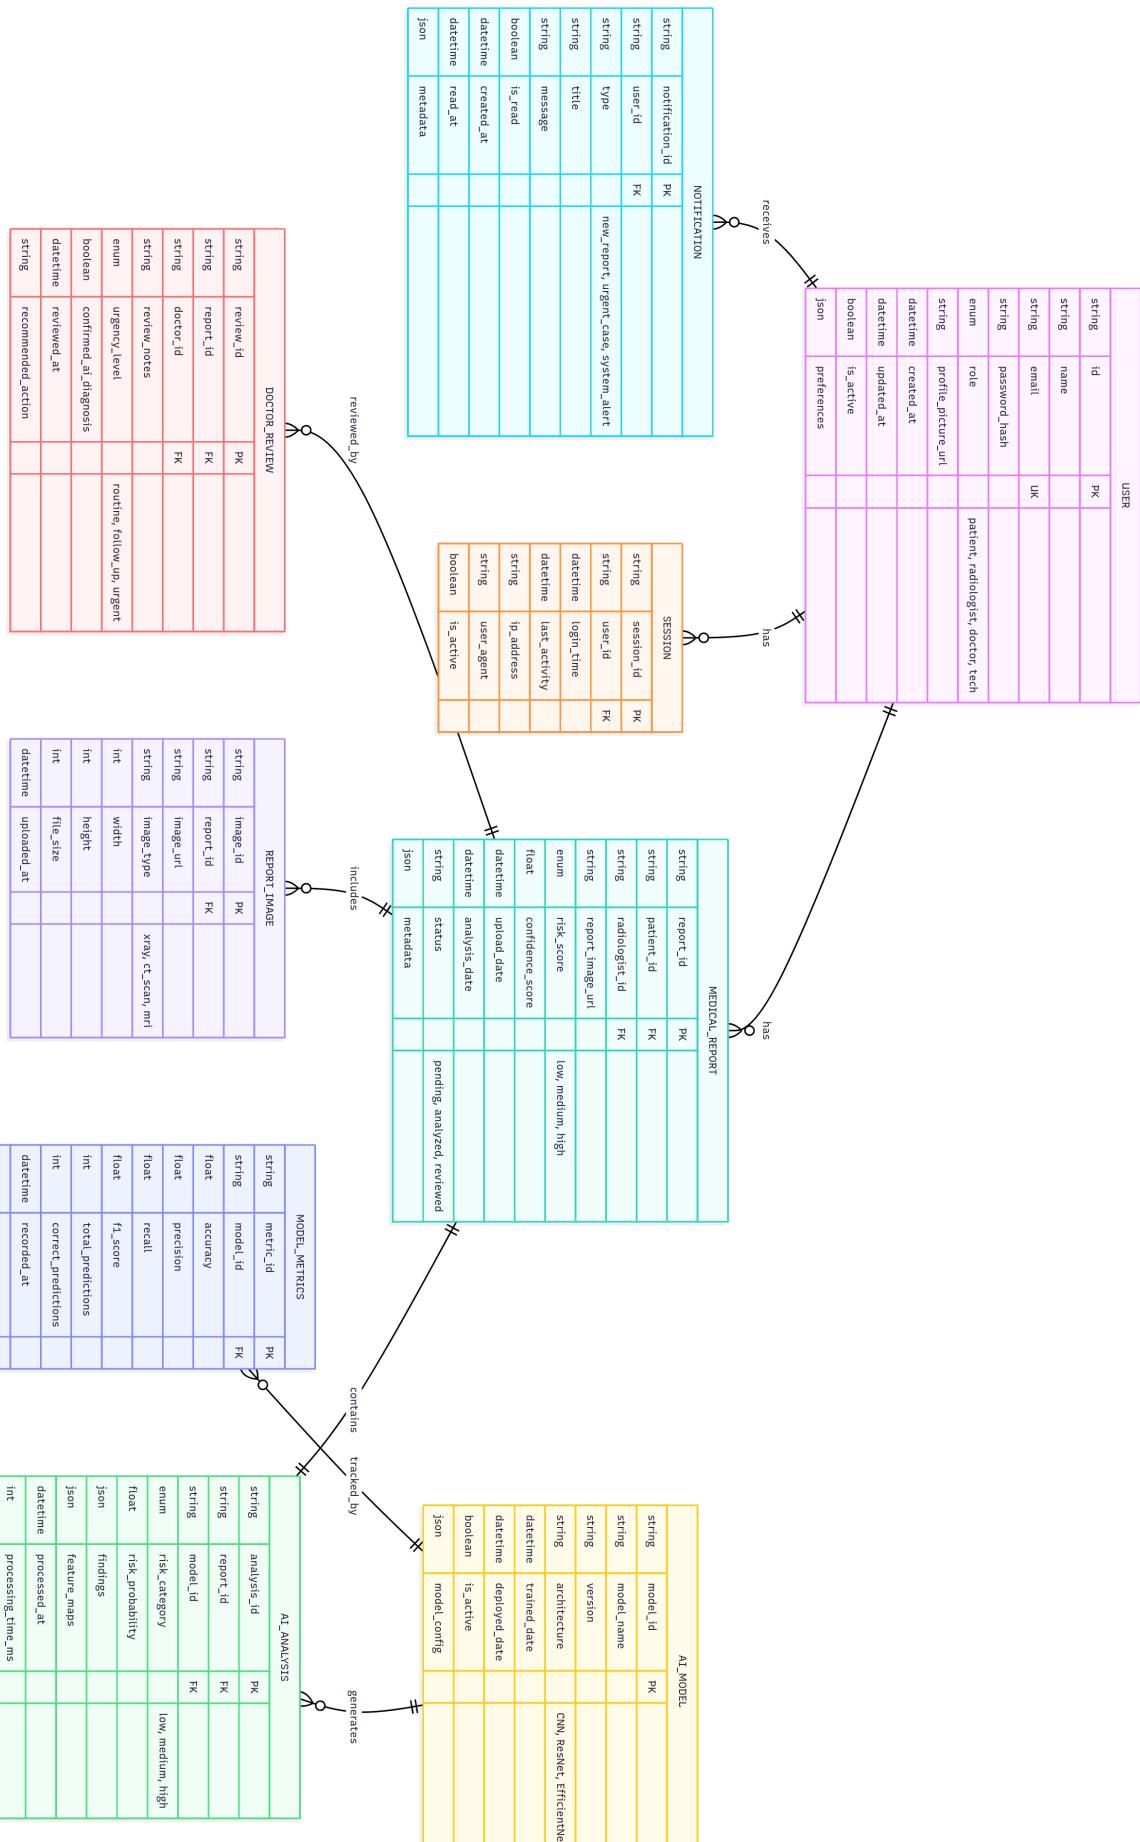


Figure 3: Entity - Relation Diagram

3.2 UML Diagram

3.2.1 User & Authentication Core

This group defines the different types of users and how they are authenticated.

- **User (Class):** This is the **base class** for all users. It contains common attributes like `id`, `name`, `email`, `role`, and `password`.
- **Radiologist & Doctor (Classes):** These classes **inherit** from **User**. This means a **Radiologist** is a **User** and a **Doctor** is a **User**.
 - Radiologist has a **specialization** and a `reviewReport()` method.
 - Doctor has a **department** and a `viewReport()` method.
- **Session (Class):** Represents a user's login session. It is linked to the **User** class (a **User** can have multiple **Sessions**).
- **JwtAuthenticationService (Class):** A service class responsible for handling the logic of `login()` and `logout()`, by creating and validating JSON Web Tokens (JWTs) for the **User**.

3.2.2 Medical Workflow Core

- **MedicalReport (Class):** A central class that holds information about a single report (`reportId`, `patientId`, `status`).
 - It is linked to a **Radiologist** (who reviews it).
 - It "has" (Aggregation) one or more **ReportImage** objects.
 - It is linked to its corresponding **AIAnalysis**.
 - It is linked to the final **DoctorReview**.
- **ReportImage (Class):** A simple class that holds metadata for an image, such as its `imageId` and `url`.
- **DoctorReview (Class):** Stores the human doctor's final notes (`reviewId`, `notes`, `urgencyLevel`), linking back to the **MedicalReport** and the **Doctor** who wrote it.

3.2.3 AI & Analysis Pipeline

This group details the classes that manage the AI models and their results.

- **AIModel (Class):** A class that represents a registered AI model, with properties like `modelId`, `name`, `version`, and `architecture`.
- **AIAnalysis (Class):** Stores the *output* of an AI model run for a specific report. It links the **MedicalReport** to the **AIModel** that was used and stores the results (`riskScore`, `findings`, `heatmapUrl`).
- **ModelMetrics (Class):** Tracks the performance of an **AIModel**, storing metrics like `accuracy`, `precision`, `recall`, etc.

3.2.4 AI Processing Sub-system

This is a detailed breakdown of the components that make up the "AI/ML Service".

- **ImageProcessor (Class):** This appears to be a *controller* or *orchestrator* class. It has an "Aggregation" relationship with the other services, meaning it *uses* them to perform a complex task. Its `processImage()` method calls the other services in sequence.
- **ImagePreprocessing (Class):** A service responsible for preparing images for the AI (e.g., `normalize()`, `resize()`).
- **RiskAssessmentEngine (Class):** The core AI logic. This class (which is linked to **AIModel**) takes the preprocessed image and `calculateRisk()` to generate a score.
- **ExplainabilityService (Class):** This service takes the AI's output and `generateHeatmap()` to create the visual explanation for the doctor.

3.2.5 Utility & Service Classes

- **Database (Class):** Represents the connection to the database. It appears to be a **Singleton** pattern (judging by `getInstance()`), which means there is only one instance of this class for the whole application.
- **NotificationService (Class):** A service responsible for sending notifications. Its `sendNotification()` method creates a **Notification** object and saves it.
- **Notification (Class):** This is the *data model* (or "entity") for a notification, holding the `message`, `userId`, `status`, etc.

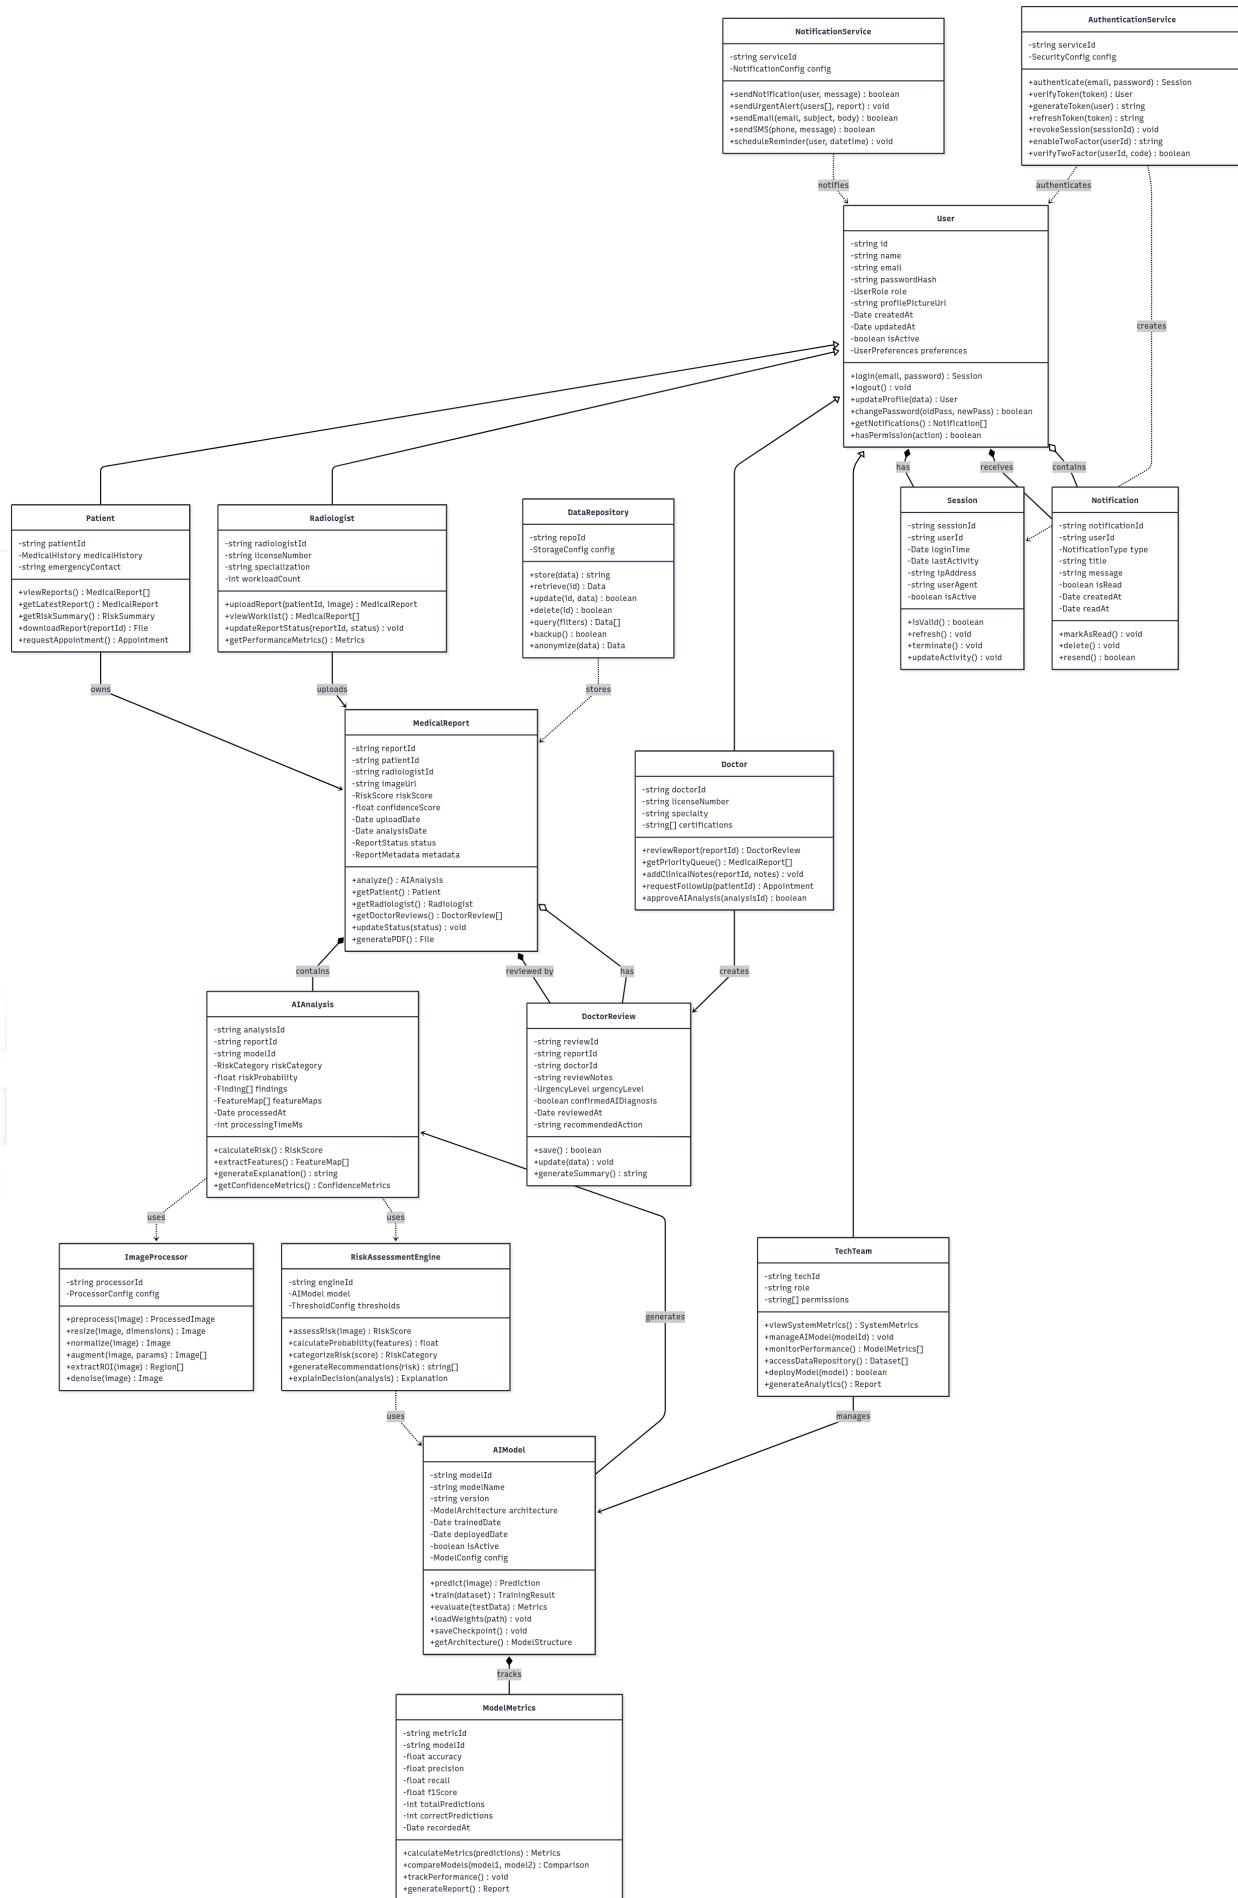


Figure 4: UML Diagram
20

4 USER FLOWS

4.0.1 Initiation and Data Ingestion

The entire process is initiated when a **Patient Report** becomes available. This report is the primary input, which is then fed directly into the **Medical AI System** to begin the automated analysis.

4.0.2 AI Analysis Pipeline

Once ingested, the report goes through a three-stage AI pipeline:

- **OCR Processing:** The system first applies Optical Character Recognition (OCR) to perform **Text Extraction**. This step converts any scanned images or non-digital text into machine-readable text.
- **Text Processing:** The extracted text is passed to a **Multi-Model NLP** (Natural Language Processing) engine. This stage understands the medical language, context, and key findings within the report.
- **Risk Assessment:** Using the features from the NLP stage, an **Ensemble Prediction** model calculates a final risk score or assessment. This is the primary analytical output of the AI.

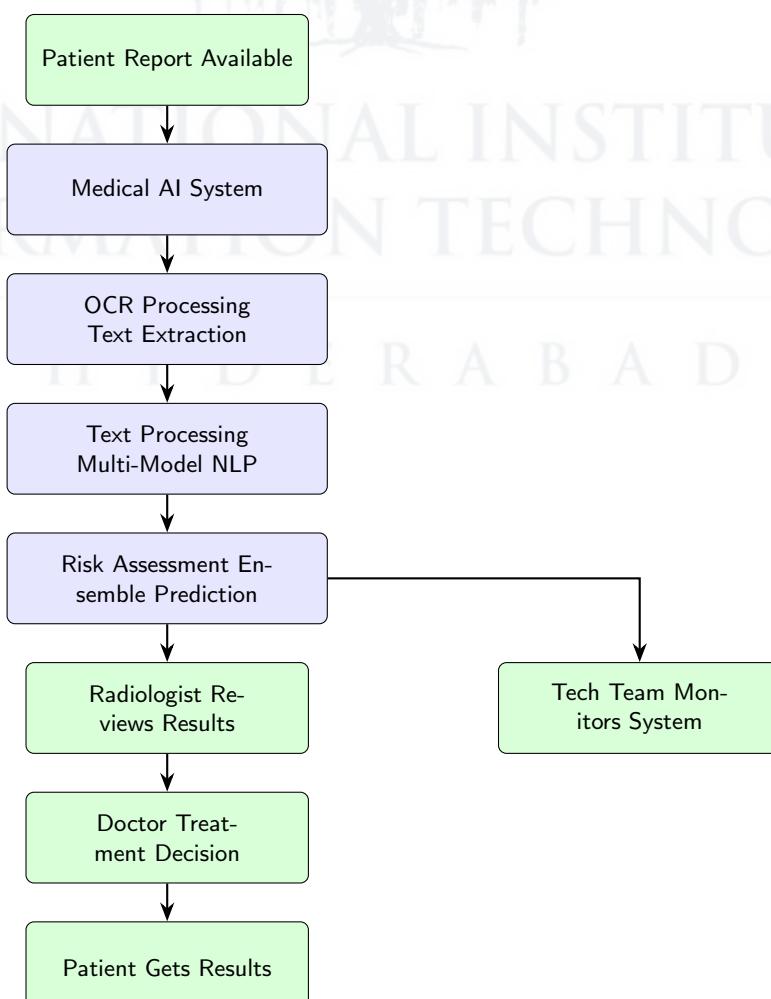
4.0.3 Workflow Branching

The output from the risk assessment triggers two parallel workflows:

- **Clinical Review (Main Path):** The results are sent to a **Radiologist** who reviews and validates the AI's findings. This human-in-the-loop step ensures clinical accuracy.
- **Technical Monitoring (Side Path):** Simultaneously, the **Tech Team** monitors the system's performance, ensuring the AI is functioning correctly, and alerts are being processed.

4.0.4 Treatment

Following the main clinical path, the radiologist's verified results are passed to the **Doctor**. The doctor uses this complete information to make a final **Treatment Decision**. The workflow concludes when the **Patient Gets Results** based on this decision.



4.1 Sequence Diagram

4.1.1 Patient Login and View Reports

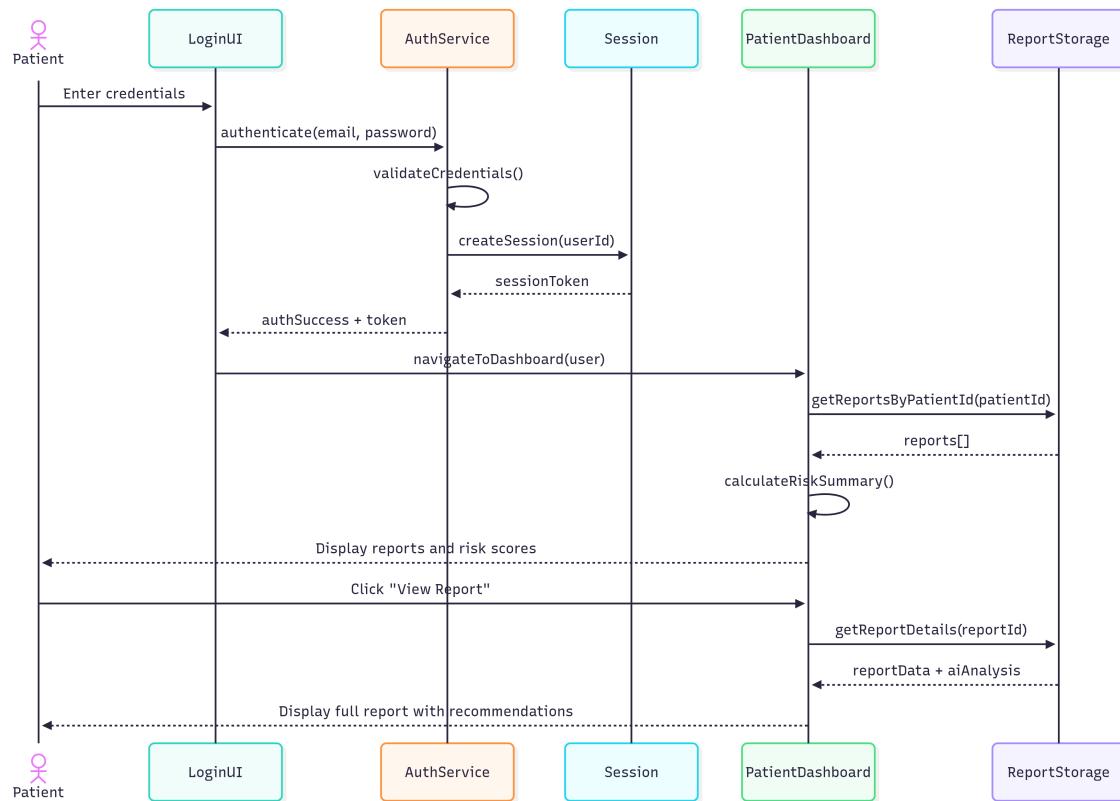


Figure 5: Patient Login and View Reports

4.1.2 Radiologist Upload X-Ray Report (AI-Powered)

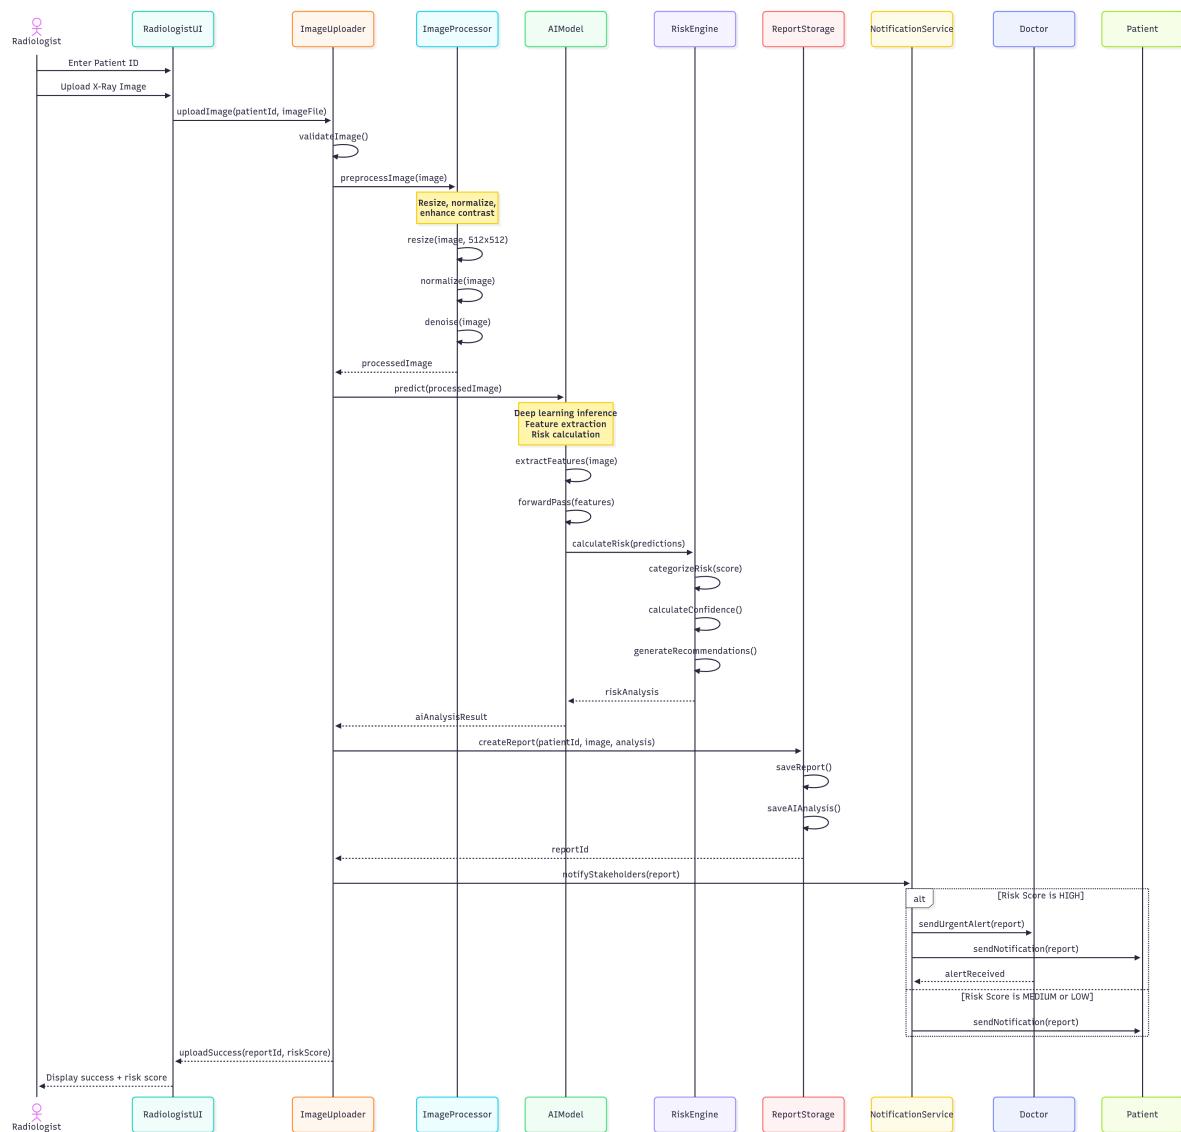


Figure 6: Radiologist Upload X-Ray Report (AI-Powered)

4.1.3 Doctor Review AI-Generated Report

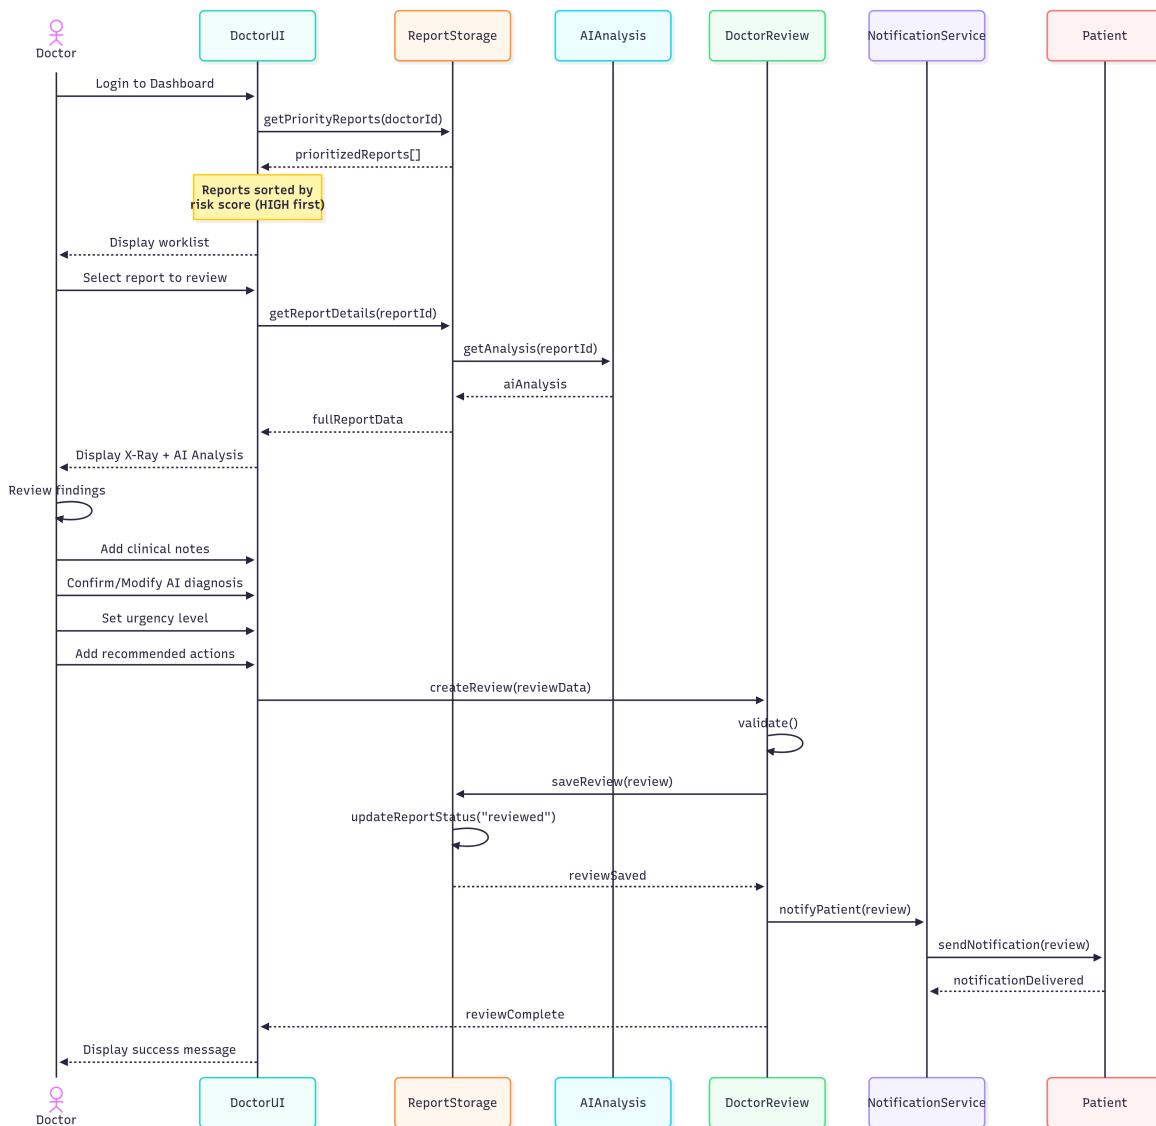


Figure 7: Doctor Review AI-Generated Report

4.1.4 Complete User Journey: From Upload to Treatment

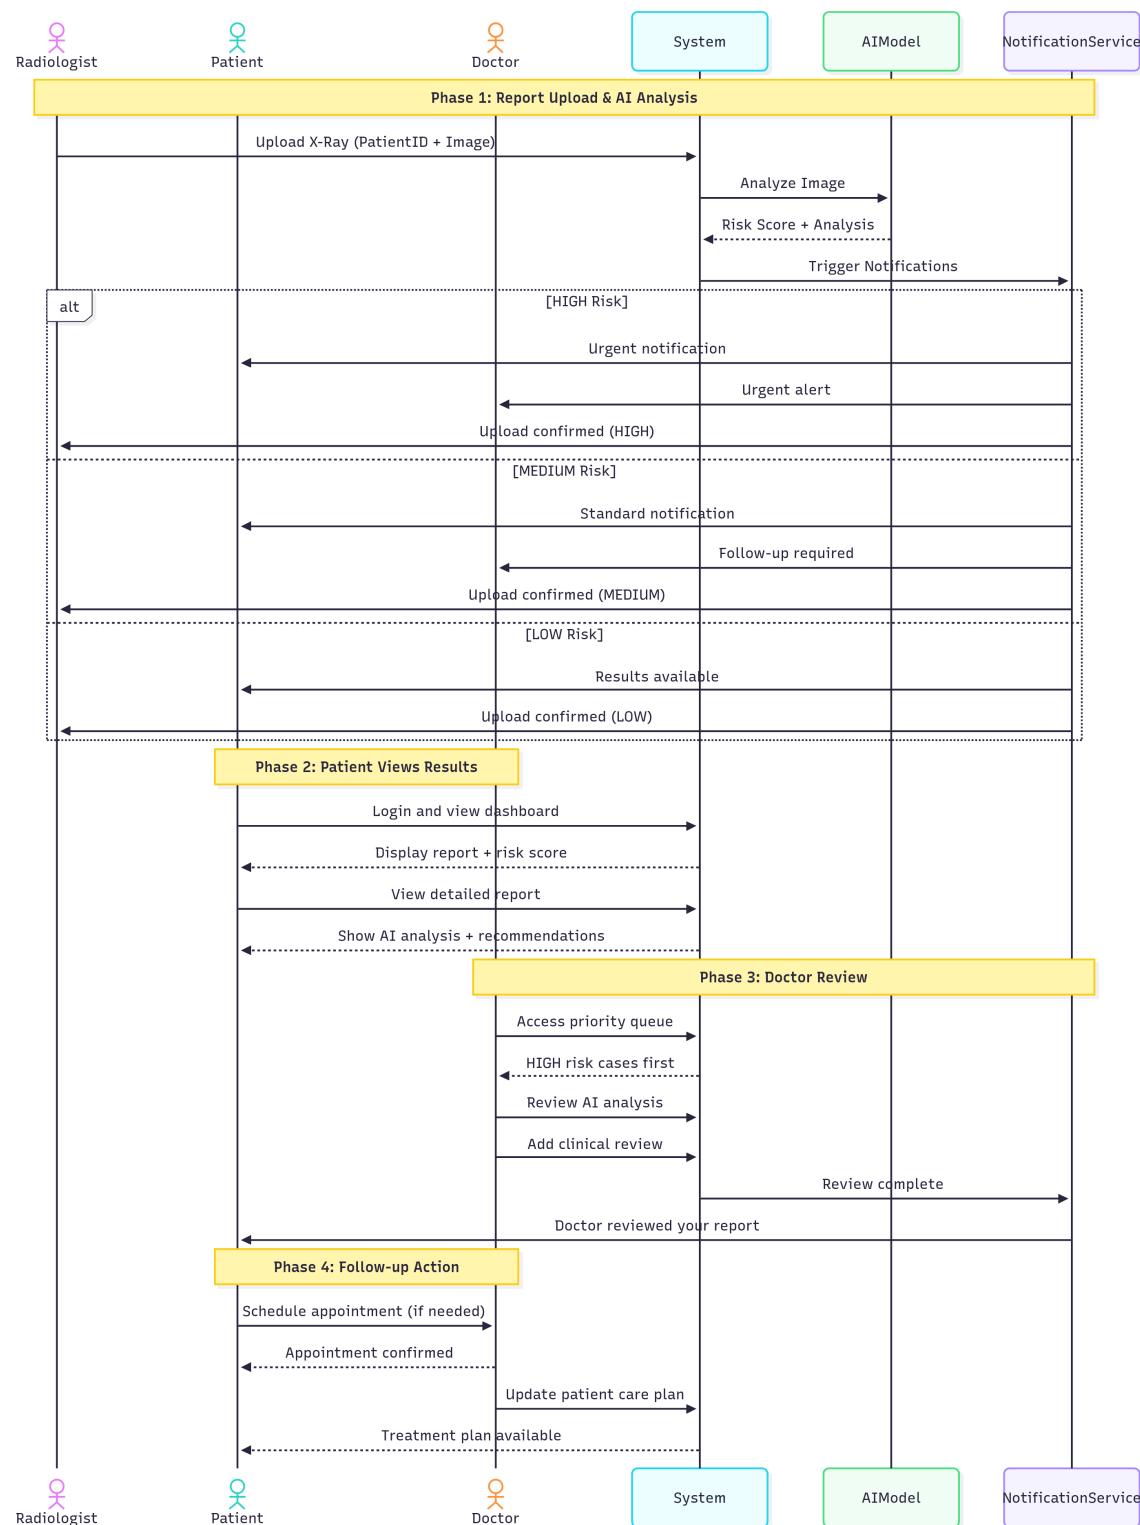


Figure 8: Complete User Journey: From Upload to Treatment

4.1.5 AI Model Prediction Pipeline

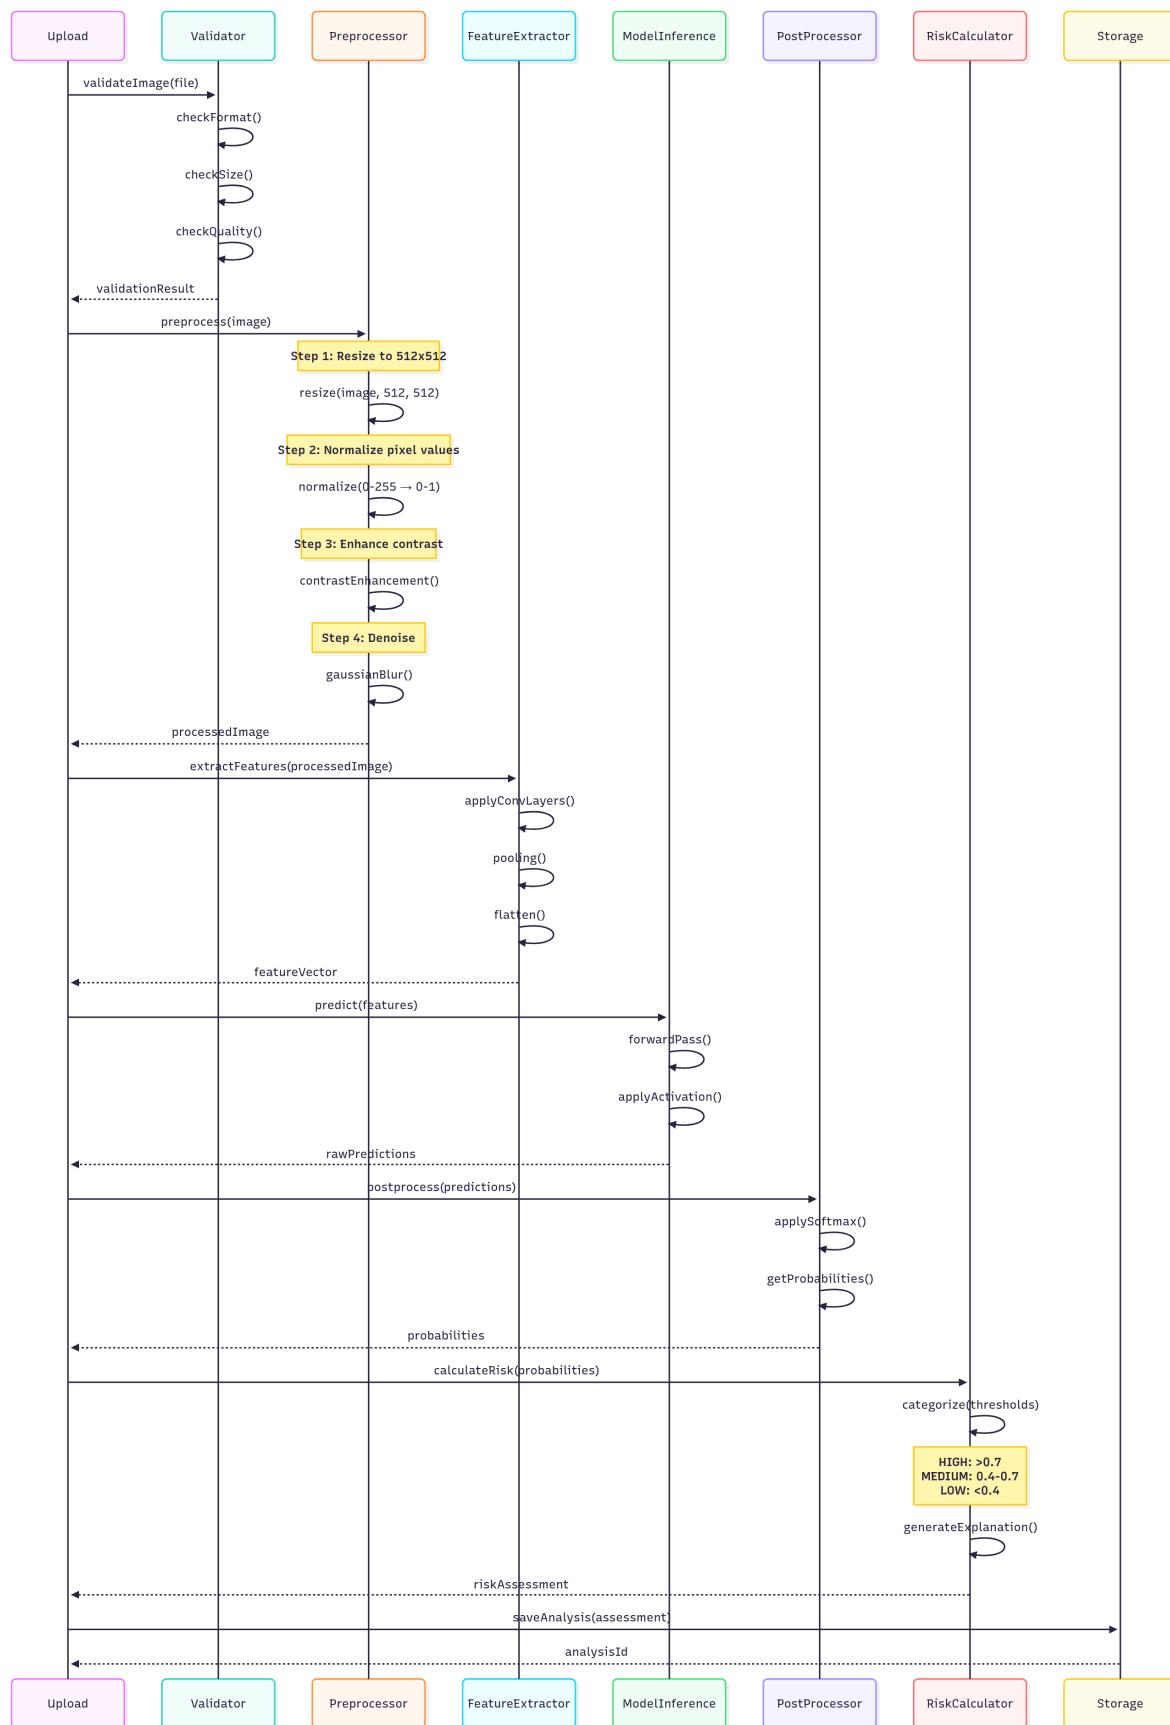


Figure 9: AI Model Prediction Pipeline

5 APIs

5.1 Endpoint Definitions

5.1.1 1. Upload & Analyze Radiology Report

This is the primary **POST** endpoint used to submit a new job to the AI. It sends a comprehensive JSON payload containing the radiology report text along with critical patient information (like age, symptoms, and history) that the model uses to provide a more accurate, contextualized analysis.

The system responds **asynchronously**, meaning it doesn't return the result immediately. Instead, it confirms the job is **accepted** and provides a **report id**, which is used to retrieve the results once they are ready.

```
{  
    "patient_id": "abc-123",  
    "radiology_report": {  
        "report_text": "Frontal chest radiograph shows bilateral lower lobe consolidation with pleural effusion...",  
        "report_date": "2025-10-10",  
        "radiologist_id": "rad-456",  
        "report_type": "chest_xray"  
    },  
    "patient_info": {  
        "age": 45,  
        "gender": "M",  
        "symptoms": ["cough", "fever", "shortness_of_breath"],  
        "medical_history": ["diabetes", "hypertension"],  
        "current_medications": ["metformin", "lisinopril"]  
    }  
}
```

Listing 1: POST /api/v1/reports/text-analyze (Request)

```
{  
    "status": "accepted",  
    "report_id": "xyz-789",  
    "estimated_time_seconds": 85,  
    "processing_stages": [  
        "text_extraction",  
        "nlp_analysis",  
        "feature_fusion",  
        "risk_calculation",  
        "explainability_generation"  
    ]  
}
```

Listing 2: POST /api/v1/reports/text-analyze (Response)

5.1.2 2. Get Text Analysis Results

This **GET** endpoint is used to retrieve the completed analysis using the **reportid** from the first call. This is the main data payload for the clinician.

- **risk assessment:** The final, top-level answer (e.g., "high" risk score).
- **text analysis:** The detailed, raw outputs from the individual NLP models (CheXpert, BioBERT).
- **ensemble predictions:** Shows how much each model contributed to the final score.
- **explainability:** The "why" behind the decision. It provides a list of **shap features** and a plain English summary.
- **recommendations:** Actionable clinical next steps based on the findings.

```
{
  "report_id": "xyz-789",
  "status": "analyzed",
  "processing_time_ms": 1285,

  "risk_assessment": {
    "risk_category": "high",
    "risk_score": 0.85,
    "confidence": 0.92,
    "urgency_level": "immediate_attention",
    "risk_factors": ["bilateral_consolidation", "pleural_effusion", "age_factor"]
  },
  "text_analysis": {
    "chexpert_findings": {
      "pneumonia": 0.89,
      "pleural_effusion": 0.82,
      "cardiomegaly": 0.23,
      "atelectasis": 0.45
    },
    "biobert_insights": {
      "clinical_context": "acute respiratory infection with complications",
      "severity_indicators": ["bilateral", "extensive", "moderate_to_severe"],
      "uncertainty_markers": [" ", "consistent_with"],
      "anatomical_locations": ["bilateral_lower_lobes", "costophrenic_angles"],
      "temporal_indicators": ["acute", "recent_onset"]
    },
    "clinical_features": {
      "symptoms_mentioned": ["consolidation", "effusion"],
      "severity_level": "moderate_to_severe",
      "laterality": "bilateral",
      "distribution": "lower_lobe_predominant"
    }
  },
  "ensemble_predictions": {
    "biobert_contribution": 0.87,
    "chexpert_contribution": 0.84,
    "xgboost_meta": 0.89,
    "gradient_boost": 0.83,
    "final_weighted": 0.85
  },
  "explainability": {
    "shap_features": {
      "most_important": [
        {"feature": "bilateral_consolidation", "importance": 0.34, "impact": "increases_risk"},
        {"feature": "pleural_effusion", "importance": 0.28, "impact": "increases_risk"},
        {"feature": "age_45", "importance": 0.15, "impact": "moderate_risk"},
        {"feature": "fever_symptoms", "importance": 0.12, "impact": "increases_risk"},
        {"feature": "diabetes_history", "importance": 0.08, "impact": "increases_risk"}
      ]
    },
    "natural_language_explanation": "The high risk score is primarily driven by bilateral lung consolidation and pleural effusion mentioned in the report, combined with patient's age and diabetic history. The bilateral nature and presence of effusion suggest a more severe respiratory condition requiring immediate attention."
  },
  "recommendations": [
    "Immediate chest CT recommended for extent assessment",
    "Start broad-spectrum antibiotic therapy",
    "Monitor oxygen saturation closely",
    "Consider ICU consultation if respiratory distress",
    "Follow-up chest X-ray in 48-72 hours"
  ],
  "quality_metrics": {
    "text_extraction_confidence": 0.98,
    "nlp_processing_quality": 0.94,
    "ensemble_agreement": 0.91
  }
}
}
```

Listing 3: GET /api/v1/reports/xyz-789 (Response)

5.1.3 3. Doctor Review and Feedback

This **POST** endpoint is the crucial **human-in-the-loop feedback mechanism**. After a doctor reviews the AI's analysis, they use this endpoint to submit their validation. It captures whether the doctor **confirmed ai diagnosis**, an **ai accuracy rating**, and the **actual diagnosis**. This feedback is essential for monitoring the AI's real-world performance and for future retraining.

```
{  
    "doctor_id": "doc-123",  
    "review_notes": "AI assessment accurate. Patient responded well to treatment.",  
    "confirmed_ai_diagnosis": true,  
    "ai_accuracy_rating": 5,  
    "actual_diagnosis": "Community-acquired pneumonia with parapneumonic effusion",  
    "treatment_plan": "IV antibiotics, chest tube placement",  
    "follow_up_required": true,  
    "follow_up_date": "2025-10-17"  
}
```

Listing 4: POST /api/v1/reports/xyz-789/review (Request)



5.2 API Signature Diagram

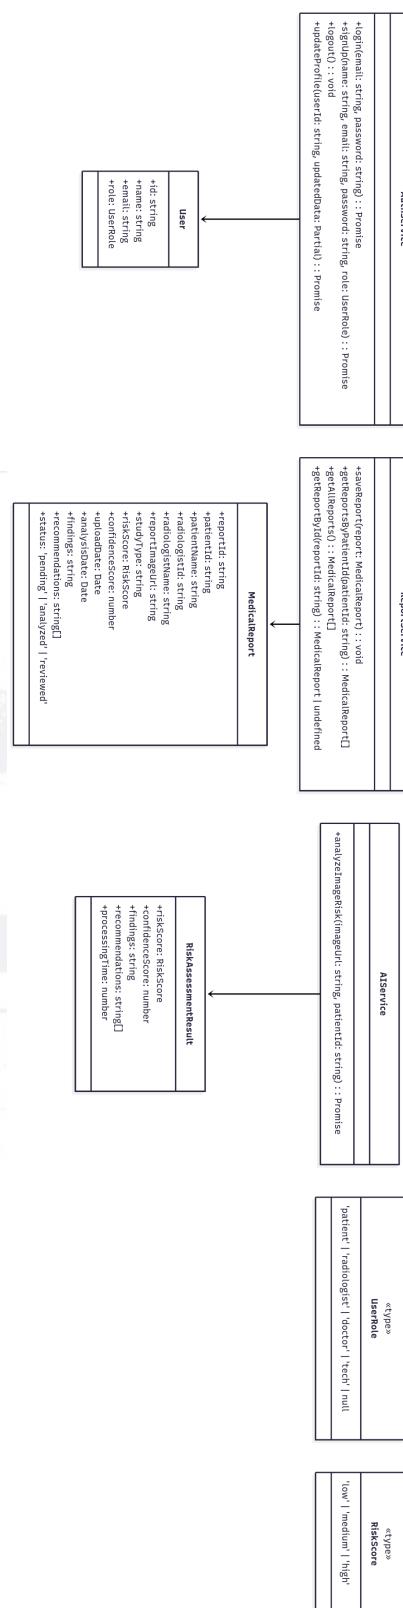


Figure 10: API Signature Diagram

6 FOLDER STRUCTURE

The project structure is as follows:

```

NNN_for_Cancer/
Code/
  backend/                      # Node.js + Python Backend
  config/                       # Configuration files
    db.js                         # MongoDB connection handler
  models/                        # Mongoose database schemas
    User.js                       # User model (patient/doctor/radiologist)
    MedicalReport.js             # Medical report with ML metadata
  routes/                        # Express API endpoints
    userRoutes.js                # User auth & management routes
    reportRoutes.js              # Report CRUD + ML integration
  utils/                         # Utility functions
    mlPipeline.js                # ML pipeline integration layer
  uploads/                       # Uploaded X-ray images storage
    .gitkeep                      # Keep directory in git
    xray-*.png                   # Uploaded image files
  Dockerfile                     # Backend container definition
  .dockerignore                  # Docker build exclusions
  .env                           # Environment variables
  .gitignore                     # Git exclusions
  server.js                      # Express server entry point
  package.json                   # Node.js dependencies
  package-lock.json              # Locked dependency versions
  requirements.txt               # Python dependencies
  risk_model.py                 # Main ML pipeline (900+ lines)
  train_xgboost_model.py        # XGBoost model training script
  xgboost_risk_model.pkl        # Trained XGBoost model
  feature_scaler.pkl            # Feature normalization scaler
  mock_model.ipynb               # Jupyter notebook for testing
  sample_report_high_risk.png   # Sample test image (high risk)
  sample_report_medium_risk.png # Sample test image (medium risk)
  x-ray-test.png                 # Sample test image

frontend/                      # React + Vite Frontend
  public/                        # Static assets
    vite.svg                      # Vite logo
  src/                           # React source code
    assets/                       # Images, icons, fonts
    components/                   # React components (26 files)
      Login.jsx                   # Authentication screen
      Login.css                   # Login styles
      Home.jsx                    # Home page component
      Home.css                    # Home page styles
      PatientDashboard.jsx       # Patient dashboard view
      PatientDashboard.css        # Patient dashboard styles
      PatientReportView.jsx      # Patient report details
      PatientReportView.css       # Patient report styles
      PatientFAQ.jsx             # Patient FAQ page
      PatientFAQ.css              # FAQ styles
      DoctorDashboard.jsx        # Doctor dashboard view
      DoctorDashboard.css         # Doctor dashboard styles
      DoctorReportView.jsx       # Doctor report review page
      DoctorReportView.css        # Doctor report styles
      RadiologistWorklist.jsx    # Radiologist work queue

```

```

RadiologistWorklist.css # Worklist styles
RadiologistUpload.jsx # Image upload interface
RadiologistUpload.css # Upload styles
T RadiologistReportInterface.jsx # Report view
RadiologistReportInterface.css # Report view styles
RadiologistArchived.jsx # Archived reports
RadiologistArchived.css # Archive styles
Reports.jsx          # General reports component
Reports.css          # Reports styles
Users.jsx            # User management component
Users.css            # User management styles
services/
    api.js           # API communication layer
    # Axios API client
App.jsx              # Main app component with routing
App.css              # Global app styles
index.css            # Root CSS styles
main.jsx             # React entry point
Dockerfile           # Frontend container (multi-stage)
.dockerignore        # Docker build exclusions
.env                 # Frontend environment variables
.gitignore           # Git exclusions
nginx.conf          # Nginx web server configuration
package.json         # Frontend dependencies
package-lock.json   # Locked dependency versions
vite.config.js       # Vite build configuration
eslint.config.js    # ESLint configuration
index.html           # HTML template

Docs/
    sequence-diagrams/
        AI_Model_Prediction_Pipeline.png
        Complete_User_Journey_from_upload_to_Treatment.png
        Doctor_Review_AI-Generated_Report.png
        patient_login_and_view_reports.png
        Radiologist_upload_X-ray_Report.png
        Tech_Team_Monitor_AI_Model_Performance.png
    schema.sql           # Initial database schema
    Architecture_design.png
    Entity-relation.png
    API_signature.png
    UML.png              # System architecture diagram
    # ER diagram
    # API signature diagram
    # UML diagrams
    api.pdf              # API documentation
    final_report.pdf     # Complete project report

docker-compose.yml
.env
deploy.sh
stop.sh
backup.sh
check-requirements.sh
PROJECT_SETUP.md
START_HERE.md
DEPLOYMENT_GUIDE.md
DEPLOYMENT_SUMMARY.md
DOCKER_README.md
    # Docker orchestration file
    # Root environment variables
    # Deployment automation script
    # Stop services script
    # Database backup script
    # System requirements checker
    # This file - complete setup guide
    # Quick start guide
    # Comprehensive deployment guide
    # Deployment overview
    # Docker commands reference

```

7 SETUP INSTRUCTIONS

7.1 Quick Setup

Quick Setup (3 Steps)

Prerequisites

- OS: Fedora Linux (or any Linux with Docker support)
- RAM: 4GB minimum (8GB recommended)
- Disk: 10GB free space
- Ports: 80, 5000, 27017 available

Step 1: Install Docker

```
# Install Docker on Fedora
sudo dnf -y install dnf-plugins-core
sudo dnf config-manager --add-repo https://download.docker.com/linux/fedora/docker-ce.repo
sudo dnf install -y docker-ce docker-ce-cli containerd.io docker-compose-plugin

# Start Docker
sudo systemctl start docker
sudo systemctl enable docker

# Add user to docker group
sudo usermod -aG docker $USER

# Apply group changes (or log out and back in)
newgrp docker

# Verify installation
docker --version
docker compose version
```

Step 2: Deploy Application

```
cd ~/NNN_for_Cancer

# Check system requirements
./check-requirements.sh

# Deploy with one command
./deploy.sh
```

Step 3: Access Application

Open browser:

- Frontend: <http://localhost>
- Backend API: <http://localhost:5000/api>

Test Credentials:

- Patient: `patient@test.com / password123`
- Doctor: `doctor@test.com / password123`
- Radiologist: `radiologist@test.com / password123`

7.2 Detailed Setup

For full details of the project setup , please refer to the repository on GitHub:

Cancer Stratification System - Github

Detailed Setup Instructions

Option A: Docker Deployment (Recommended)

1. Clone/Navigate to Project:

```
cd ~/NNN_for_Cancer
```

2. Configure Environment:

```
# Create environment file  
cp .env.example .env
```

```
# Edit if needed  
nano .env
```

3. Build & Deploy:

```
# Build all Docker images  
docker compose build
```

```
# Start all services  
docker compose up -d
```

```
# Check status  
docker compose ps
```

```
# View logs  
docker compose logs -f
```

4. Verify Deployment:

```
# Test backend  
curl http://localhost:5000/api/users
```

```
# Test frontend  
curl http://localhost
```

```
# Access MongoDB  
docker exec -it cancer-stratification-db mongosh -u admin -p adminpassword123
```

Option B: Manual Setup (Development)

Backend Setup:

```
cd Code/backend
```

```
# Install Node.js dependencies  
npm install
```

```
# Install Python dependencies  
pip3 install -r requirements.txt
```

```
# Install Tesseract OCR  
sudo dnf install tesseract tesseract-langpack-eng
```

```
# Create uploads directory  
mkdir -p uploads
```

```
# Set environment variables  
export MONGODB_URI="mongodb://localhost:27017/cancer_stratification"  
export PORT=5000
```

```
# Start backend  
npm start
```



Frontend Setup:

```
cd Code/frontend  
  
# Install dependencies  
npm install  
  
# Start development server  
npm run dev
```

Database Setup:

```
# Install MongoDB  
sudo dnf install mongodb-org  
  
# Start MongoDB  
sudo systemctl start mongod  
  
# Create database and users (use mongosh)
```

8 INDIVIDUAL CONTRIBUTIONS

- **Nilavra Ghosh** — Led the codebase setup, explored and compared possible technology stacks, management, feedback and contributed to Design documentation(Figma/ER Diagram/Sequence Diagram) and workflow structuring.
- **Ved Prakash Maurya** — Integrated the backend pipeline alongside Aviral , performed functional and regression testing, and ensured seamless interaction between backend services and frontend modules .
- **Aviral Malhotra** — Worked on the frontend development , integrated model outputs into the UI alongside Ved, coordinated the design of user flows to ensure a smooth patient–doctor interaction .
- **Bhaskar Bhatt** — Supported backend integration tasks, collaborated on data preprocessing pipelines, and contributed to system testing and debugging.
- **Akshat** — Worked on testing and integration methodology, evaluated model outputs during deployment trials, assisted in pipeline integration, and helped refine the system through iterative testing cycles.
- **Aayush** — Conducted research on the problem statement, reviewed medical literature, and helped refine the scope and framing of the project from a social innovation perspective.